

Hypothesis Testing and Confidence Intervals via Sampling Distributions generated in R

```
source("myFunctions.R") # load custom functions
```

First Approach: Make (a justified) assumption about population distr.

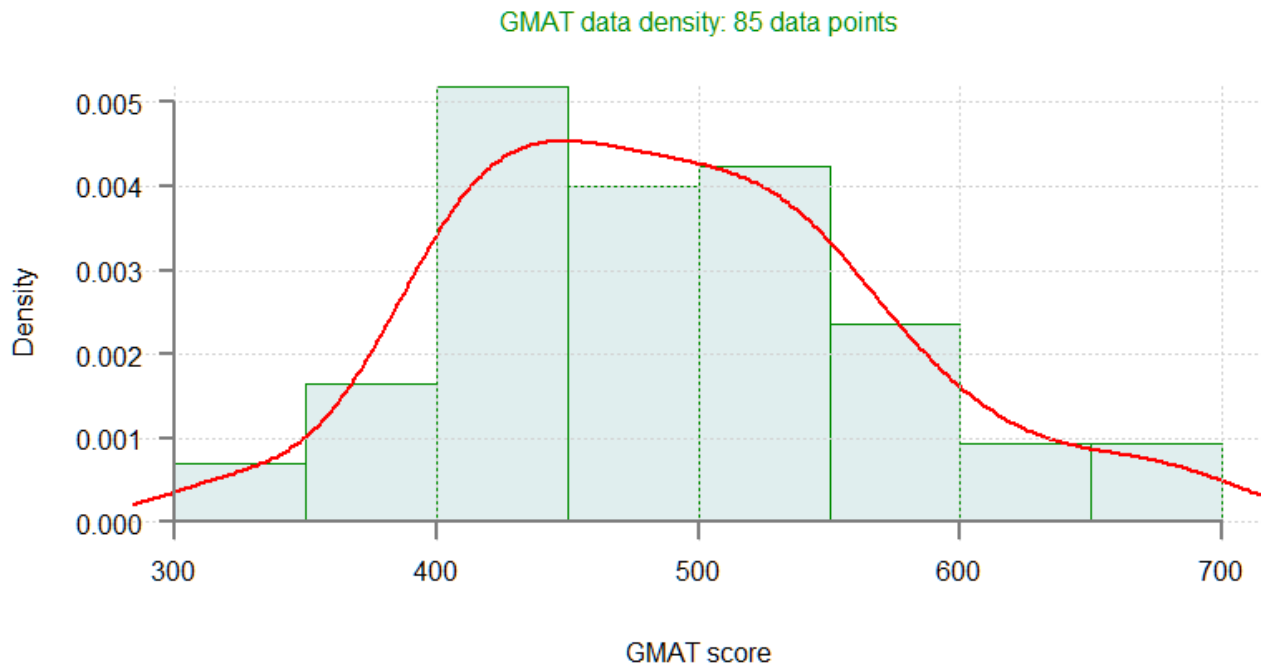
Example 1: (example from Session 8)

```
# Test the null hypothesis that the population mean of GMAT data is 510.
```

```
adm <- read.csv("Data08/Admission.csv"); head(adm)
GMAT = adm$GMAT
summary(GMAT)
```

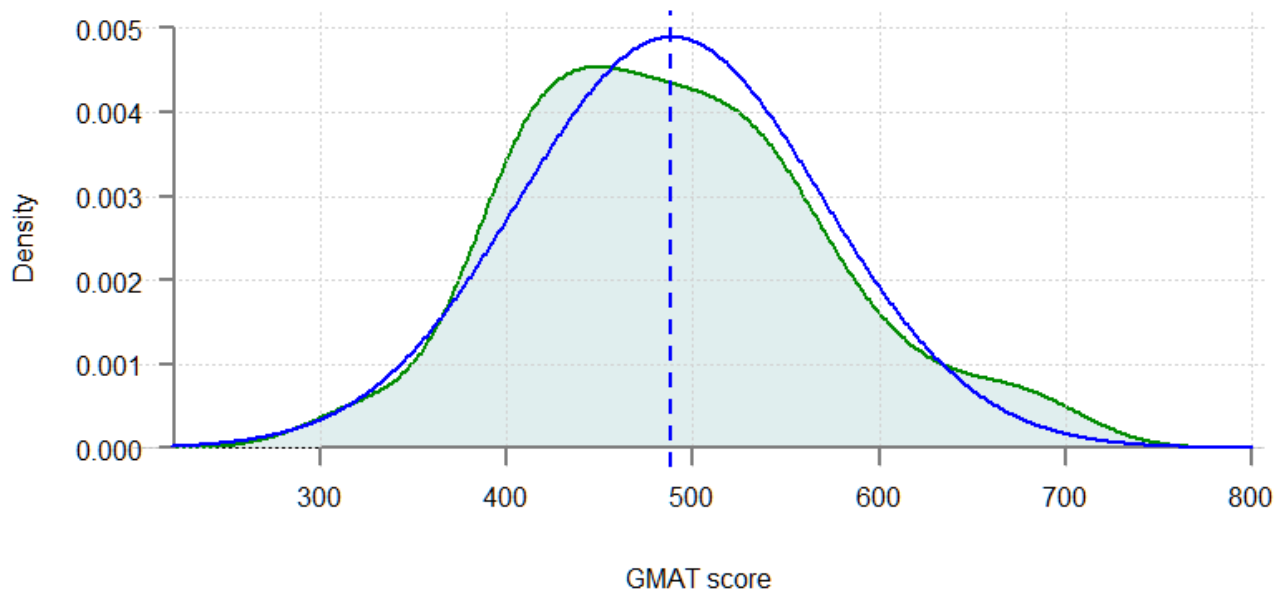
Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
313.0	425.0	482.0	488.4	538.0	693.0

```
# From Session 9: GMAT data density, sample mean, etc.
myhd(GMAT,ylim=c(0,0.005),col="azure2",border="green4",
     xlab="GMAT score",main="GMAT data density")
```



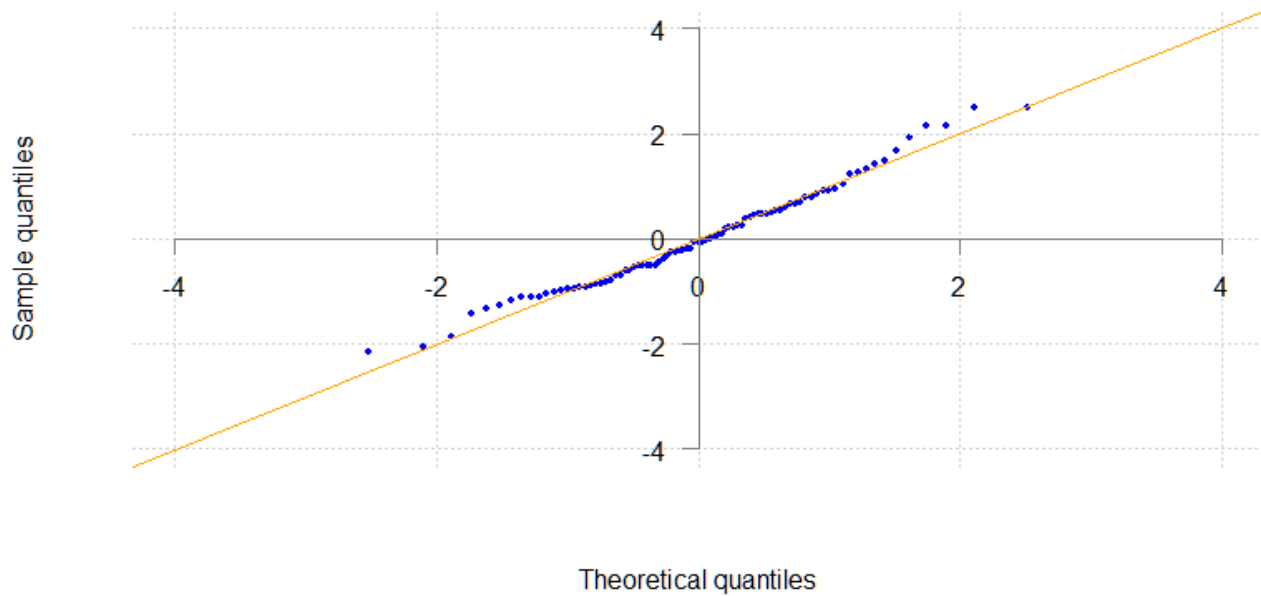
```
myed(GMAT,ylim=c(0,0.005),pcol="azure2",dcol="green4",
     xlab="GMAT score",main="GMAT data density")
sm = mean(GMAT); sm
ssd = sd(GMAT); ssd
x = seq(220,800)
lines(x,dnorm(x,mean=sm,sd=ssd),lwd=2,col="blue")
abline(v=sm,lwd=2,lty=2,col="blue")
```

GMAT data density: 85 data points



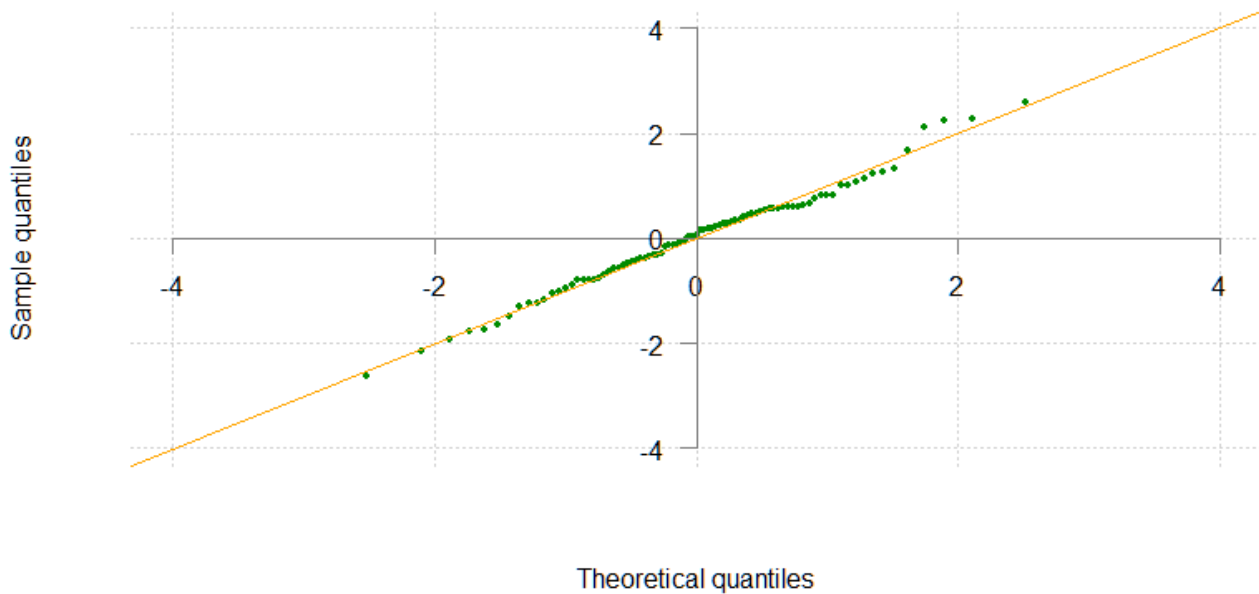
```
# Is GMAT data normally distributed? Try with a snQQplot:  
snQQplot(trv(GMAT),main="Transf. GMAT")
```

Normal Q-Q Plot: 85 Transf. GMAT points



```
# How different is it from the snQQplot of the true Normal random sample?  
snQQplot(trv(rnorm(85,sm,ssd)),col="green4",main="Transf. Normal")
```

Normal Q-Q Plot: 85 Transf. Normal points



```
# KEY: Based on observed GMAT density and the quantile plot WE ASSUME
#       the population distribution is Normal.
#       (discussed in great detail, and justified, in Session 9+)
# The best (unbiased) estimators for population mean and variance are
# Sample Mean and Sample Variance, respectively.
# Thus we assume population distribution is
#        $N(sm, ssd^2) = N(488.4, 81.5(\text{square}))$ 

# The sample size is n = 85 (there are 85 rows in the dataframe).
n = length(GMAT)

# To create one sample
rnorm(n, sm, ssd)

# To compute the metric (mean) for one sample
mean(rnorm(n, sm, ssd))

# Finally, replicate this 50 thousand times to get a
# Sampling Distribution Vector for the Population Mean
m = 50000
sdv1a = replicate(m, mean(rnorm(n, sm, ssd)))

# Plot the empirical density of the sampling distribution vector

# Note: 'Sampling Distribution' and 'Sampling Distribution Vector'
#       are synonyms. The generated vector is a 'representative'
#       of the distribution we (usually) don't know exactly.
#       The way we visualize vector 'representing' the distribution
#       is via the empirical density. Hence it is reasonable to refer to
#       the Empirical Density of Sampling Distribution Vector as
#       'Sampling Distribution' as well.

myed(sdv1a, cntFlg=F,
     main=paste("Pop.Mean Sampling Distribution from Normal sample of size", n),
     sub="assuming GMAT data mean and variance")

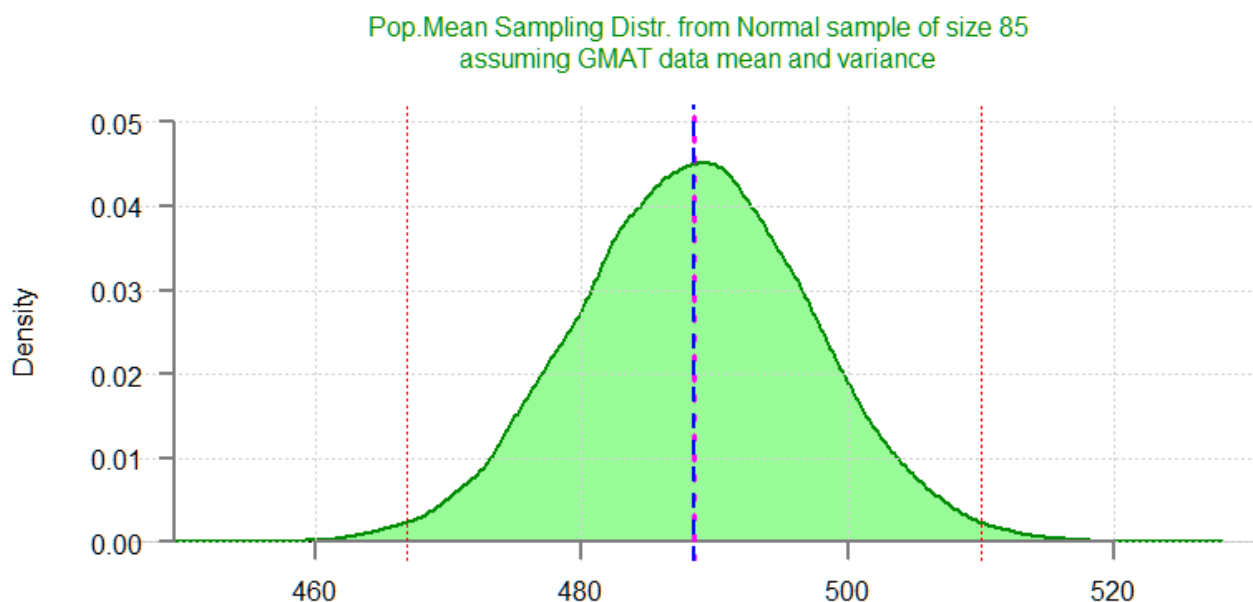
myed(sdv1a, ylim=c(0, 0.05), cntFlg=F,
     main=paste("Pop.Mean Sampling Distr. from Normal sample of size", n),
     sub="assuming GMAT data mean and variance")
```

```
# without custom functions:
# d_sdv1a = density(sdv1a)
# plot(d_sdv1a,lwd=3,col="green2",
#      main=paste("Pop.Mean Sampling Distribution from Normal sample of size",n),
#      "assuming GMAT data mean and variance")); grid()
# polygon(d_sdv1a,col="lightgreen")

# Note: in this case we know the true distribution of the sampling
#       distribution: N(488.4, 81.5/sqrt(85))
# Question: what is the "half-mass" point (median)?
abline(v=median(sdv1a),lwd=3,col="magenta",lty=3)
# Where does the sm fit in?
abline(v=sm,lwd=2,lty=2,col="blue")

# Null-hypothesis H0: GMAT population mean is 510
mu0 = 510

# How far from sm is the mu0?
abline(v=mu0,lwd=1.5,col="red",lty=3)
abline(v=sm+sm-mu0,lwd=1.5,col="red",lty=3)
```



```
# Empirical p-value (see the plot)
left_p = length(sdv1a[sdv1a<sm+sm-mu0])
right_p = length(sdv1a[sdv1a>mu0])
pval = (left_p+right_p)/length(sdv1a); pval

# In order to avoid keeping track of which of the two values,
# mu0 and sm+sm-mu0 is larger/smaller,
lowbd = min(mu0,sm+sm-mu0) # lower bound
uppbld = max(mu0,sm+sm-mu0) # upper bound

# Compute the probabilities 'beyond' these lines to obtain
# the p-values for the hypothesis
left_p = length(sdv1a[sdv1a<lowbd])
right_p = length(sdv1a[sdv1a>uppbld])
pval = (left_p+right_p)/length(sdv1a); pval
```

```
[1] 0.01534
```

```
# We reject the hypothesis at significance level 0.05.
```

```
# Since this is a test for the population mean, we can compare to built-in t-test:
```

```
t.test(GMAT,mu=mu0,conf.level = 0.95)
```

```
One Sample t-test
```

```
data: GMAT
```

```
t = -2.4375, df = 84, p-value = 0.0169
```

```
alternative hypothesis: true mean is not equal to 510
```

```
95 percent confidence interval: 470.8631 506.0310
```

```
sample estimates:
```

```
mean of x      488.4471
```

```
# Related question: What is the empirical conf.interval for 95% confidence?
```

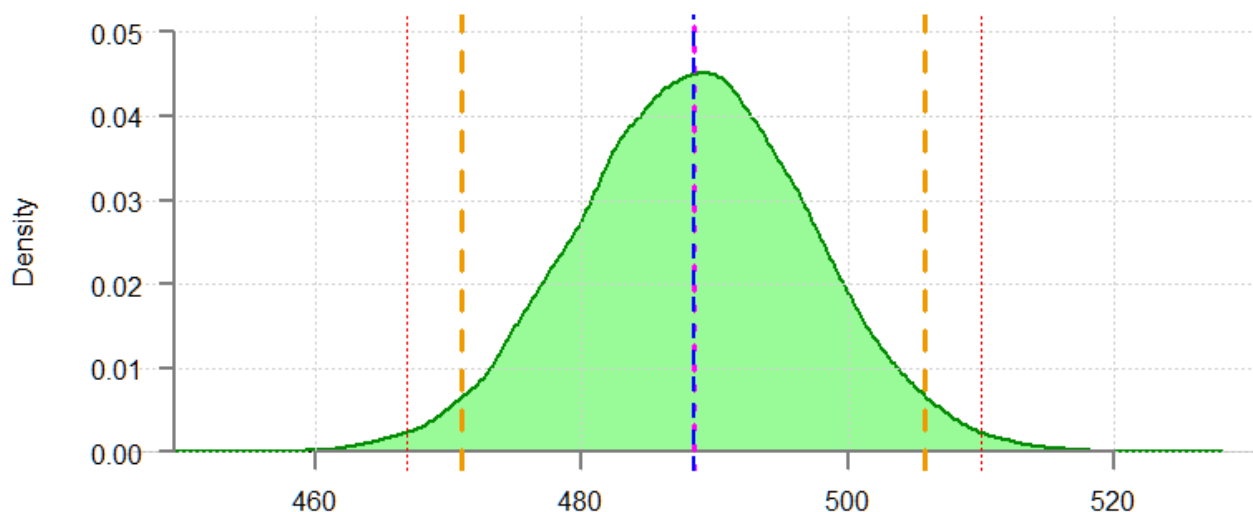
```
myeCI(sdv1a,0.95)
```

```
2.5%    97.5%
```

```
471.0071 505.7843
```

```
abline(v=myeCI(sdv1a,0.95),lwd=3,lty=2,col="orange2")
```

Pop.Mean Sampling Distr. from Normal sample of size 85
assuming GMAT data mean and variance



```
# without custom functions:
```

```
# left_ep = quantile(sdv,0.025)
```

```
# right_ep = quantile(sdv,0.975)
```

```
# CI = c(left_ep,right_ep); CI
```

```
# abline(v=CI,lwd=1.5,col="purple")
```

```
# As in the classical t-test we shall often AVOID the computation of p-values
```

```
# and use empirical confidence interval instead.
```

```
# Conclusion: Since the hypothesized value  $\mu_0 = 510$  is outside the two-sided
```

```
# empirical 95% confidence interval, we reject the null hypothesis
```

```
# in favor of the alternative: Population mean is not equal 510.
```

```
# Useful wrap (to avoid cutting and pasting same lines over and over again)
```

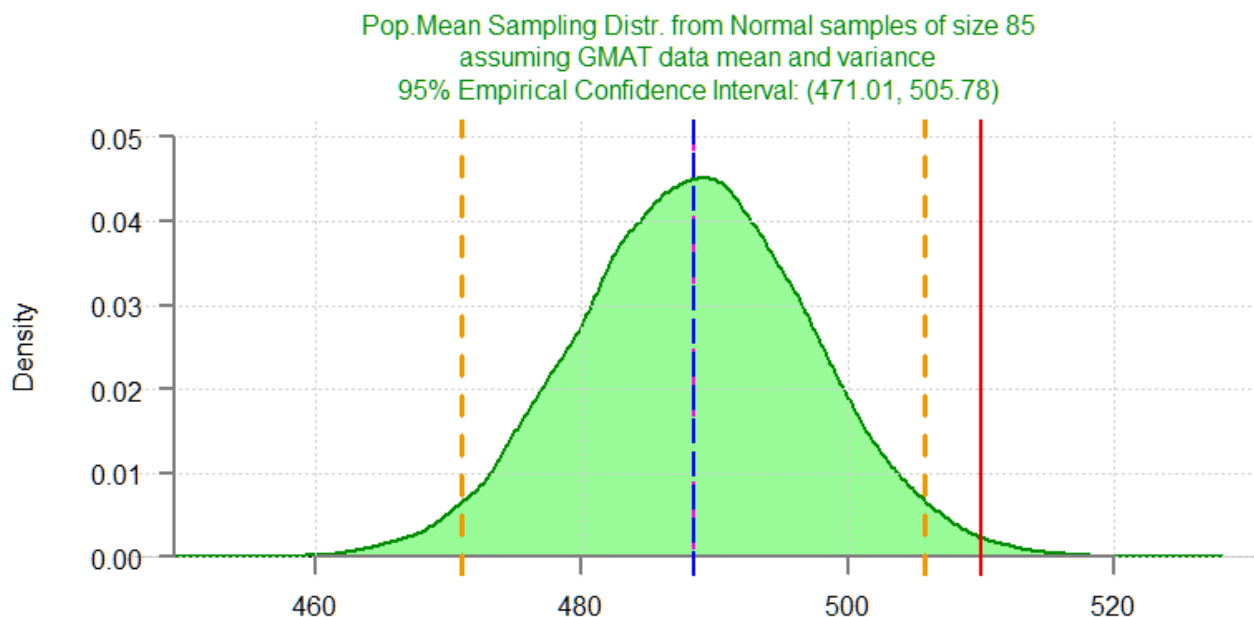
```
# Function plots the sampling distribution density and marks confidence interval
```

```
# on the plot. It also returns the confidence interval endpoints.
```

```
mysdci = function(sdv,xlim=NULL,ylim=NULL,xlab="",main=NULL,sub=NULL,cntFlg=F,
                  pcol="palegreen",dlwd=2,dcol="green4",axes=T,conf=0.95,rnd=2)
{
  CI = myeCI(sdv,conf)
  ssub = paste(100*conf,"% Empirical Confidence Interval: (",
               round(CI[1],rnd),", ",round(CI[2],rnd),")",sep="")
  if (!is.null(sub)) sub = c(sub,ssub) else sub = ssub
  myed(sdv,xlim=xlim,ylim=ylim,xlab=xlab,main=main,sub=sub,cntFlg=cntFlg,
       pcol=pcol,dlwd=dlwd,dcol=dcol,axes=axes)
  abline(v=median(sdv),lwd=2,col="magenta",lty=2)
  abline(v=CI,lwd=3,lty=2,col="orange2")
  return(CI)
}
```

Apply it to 'sdv1a' sampling distribution vector:

```
mysdci(sdv1a,ylim=c(0,0.05),
       main=paste("Pop.Mean Sampling Distr. from Normal samples of size",n),
       sub="assuming GMAT data mean and variance")
abline(v=sm,lwd=2,col="blue",lty=5)      # GMAT data sample mean (average)
abline(v=mu0,lwd=2,col="red2")           # hypothesized value:  $\mu_0 = 510$ 
```



Note: In classical theory CI for population mean depends only on
the sample and the confidence level.
In this case the empirical CI depends on the generated sdv.
How much of the variation should we see?
We can certainly repeat CI calculations say $M = 100$ times.
But we must reduce m to avoid long execution times

```
m = 1000; M = 100
CIsdv1a = rpl(M,myeCI(rpl(m,mean(rnorm(n,sm,ssd)))))

par(mfrow=c(2,1))
myhd(CIsdv1a[1,],breaks=10,cntFlg=F,main="95% CI's left end-points")
myhd(CIsdv1a[2,],breaks=10,cntFlg=F,main="95% CI's right end-points",
     col="lightblue",border="blue2",dcol="blue")
par(mfrow=c(1,1))
```

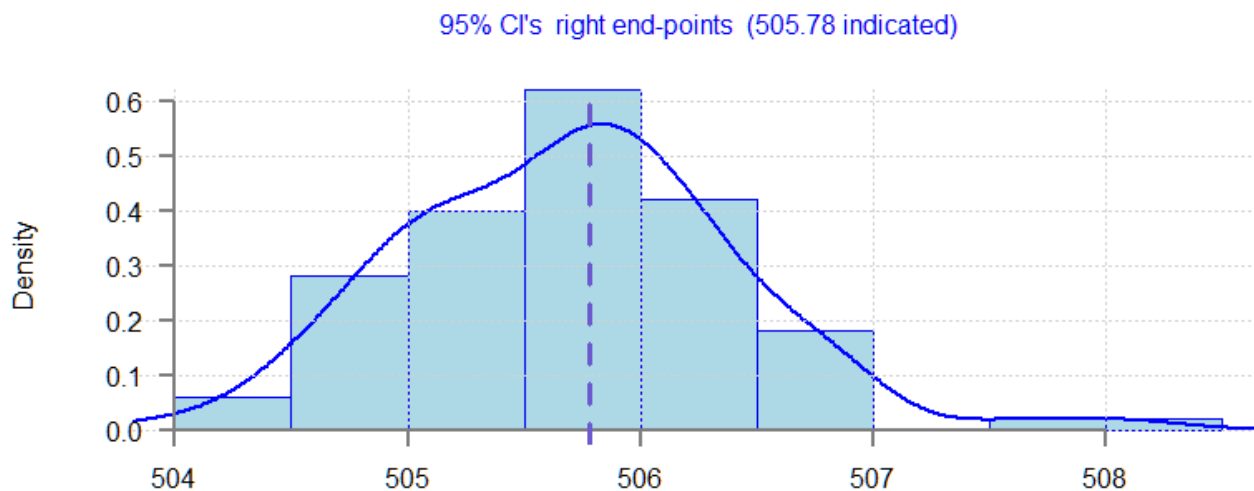
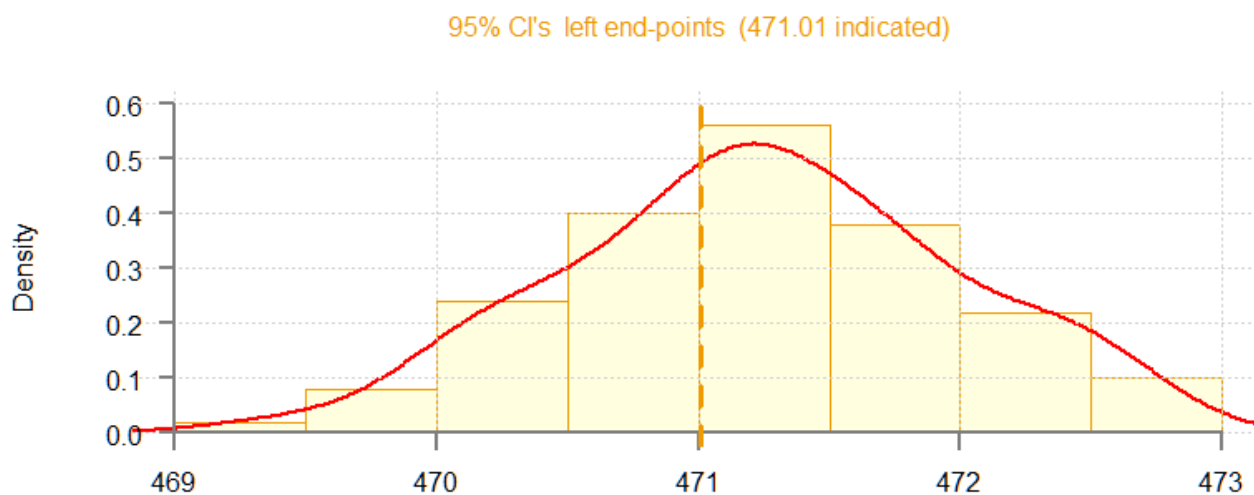
wrap it as an auxiliary function:

```

myCIsdv = function(CIsdv,marker=NULL,ylim=NULL,conf=0.95,rnd=2)
{
  par(mfrow=c(2,1))
  b = ceiling(sqrt(length(CIsdv)/2))
  main = paste(100*conf,"% CI's ",sep=""); sub=NULL
  if (!is.null(marker))
    sub = paste(" (",round(marker[1],rnd)," indicated)",sep="")
  myhd(CIsdv[1,],breaks=b,ylim=ylim,cntFlg=F,main=paste(main,"left end-points",sub))
  abline(v=marker[1],lwd=3,lty=2,col="orange2")
  if (!is.null(marker))
    sub = paste(" (",round(marker[2],rnd)," indicated)",sep="")
  myhd(CIsdv[2,],breaks=b,ylim=ylim,cntFlg=F,main=paste(main,"right end-points",sub),
    col="lightblue",border="blue2",dcol="blue")
  abline(v=marker[2],lwd=3,lty=2,col="slateblue")
  par(mfrow=c(1,1))
}

myCIsdv(CIsdv1a,myeCI(sdv1a))
myCIsdv(CIsdv1a,myeCI(sdv1a),ylim=c(0,0.6))

```



Example 2:

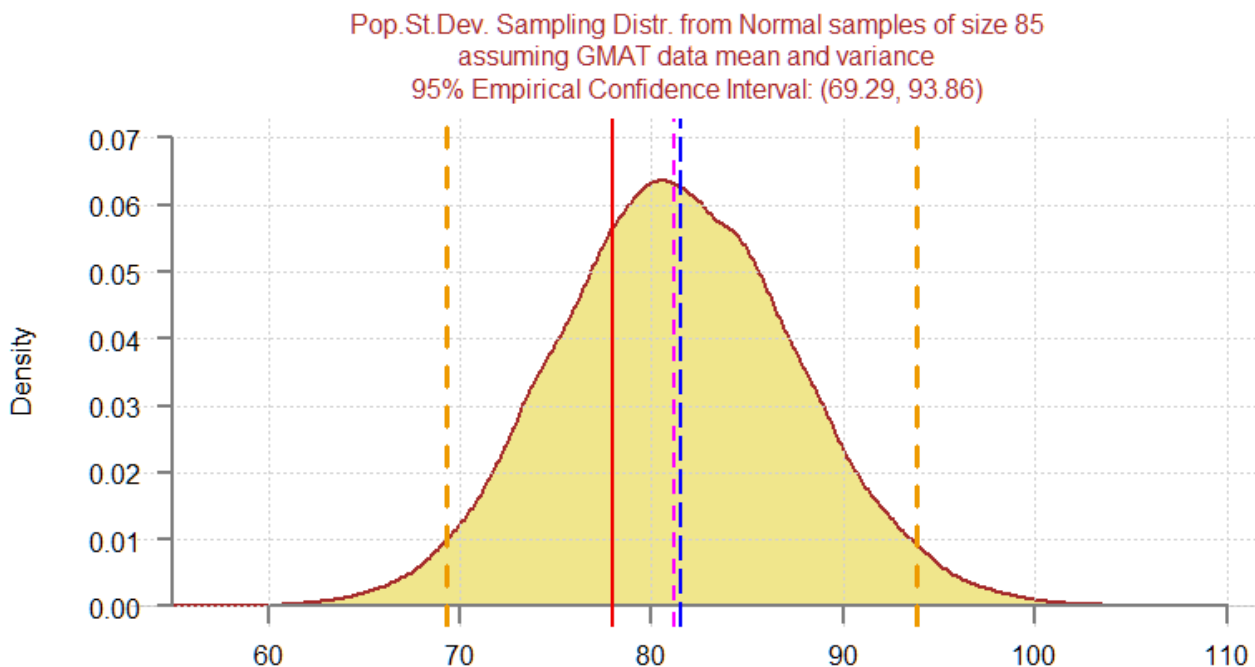
```
# Empirical Confidence Interval for population standard deviation
# "Null-hypothesis: population standard deviation is 78"

# Lineup as before: Sample size, mean, and standard deviation of GMAT data
# n = length(GMAT); n
# sm = mean(GMAT); sm
# ssd = sd(GMAT); ssd

# Hypothesized Population Standard Deviation value
sd0 = 78

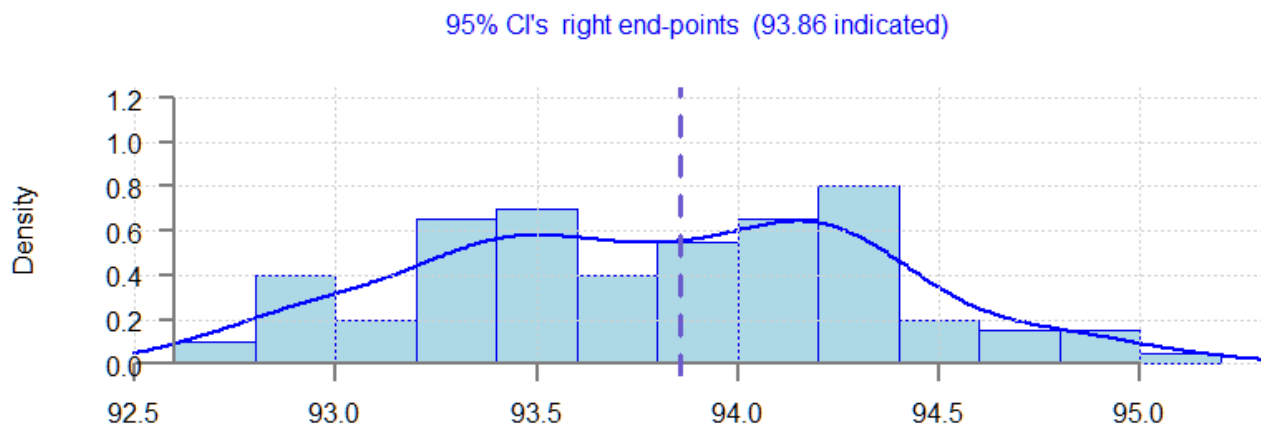
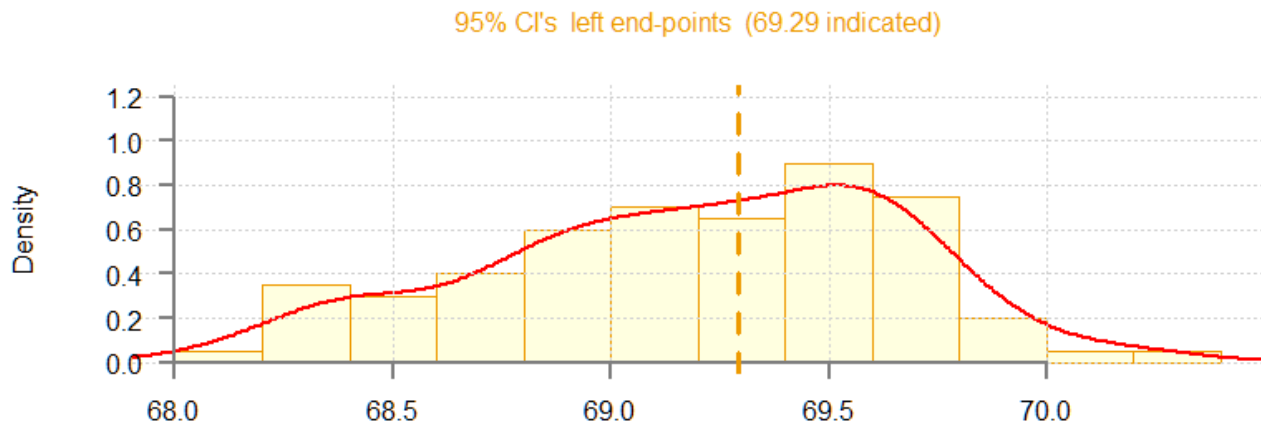
# "Metric" is the sample standard deviation: we need to replicate it many times
sd(rnorm(n,sm,ssd))

# Get the Sample Standard Deviation of 50,000 replicates
# (stored as sampling distribution vector) and plot its density
m = 50000
sdv2a = replicate(m,sd(rnorm(n,sm,ssd)))
mysdci(sdv2a,ylim=c(0,0.07),pcol="khaki",dcol="brown",
       main=paste("Pop.St.Dev. Sampling Distr. from Normal samples of size",n),
       sub="assuming GMAT data mean and variance")
abline(v=sd0,lwd=2,col="red2")      # hypothesized value: sd0 = 78
abline(v=ssd,lwd=2,col="blue",lty=5) # sample sd of GMAT data
```



```
# Test conclusion: We accept the null-hypothesis at significance level 0.05.
```

```
# CI variation:
m = 1000; M = 100
CIsdv2a = rpl(M,myeCI(rpl(m,sd(rnorm(n,sm,ssd)))))
myCIsdv(CIsdv2a,myeCI(sdv2a))
myCIsdv(CIsdv2a,myeCI(sdv2a),ylim=c(0,1.2))
```

Example 3:

```
# Empirical Confidence Interval for the population 75th percentile
# "Null-hypothesis: Population 75th percentile is 600"

qnt = 0.75      # 75th percentile = 0.75 quantile
# qnt = 0.99    # 99th percentile
# qnt = 1       # maximum

# Sample size, mean, standard deviation, and the desired percentile of GMAT data
# n = length(GMAT); n
# sm = mean(GMAT); sm
# ssd = sd(GMAT); ssd
sqnt = quantile(GMAT,qnt); sqnt

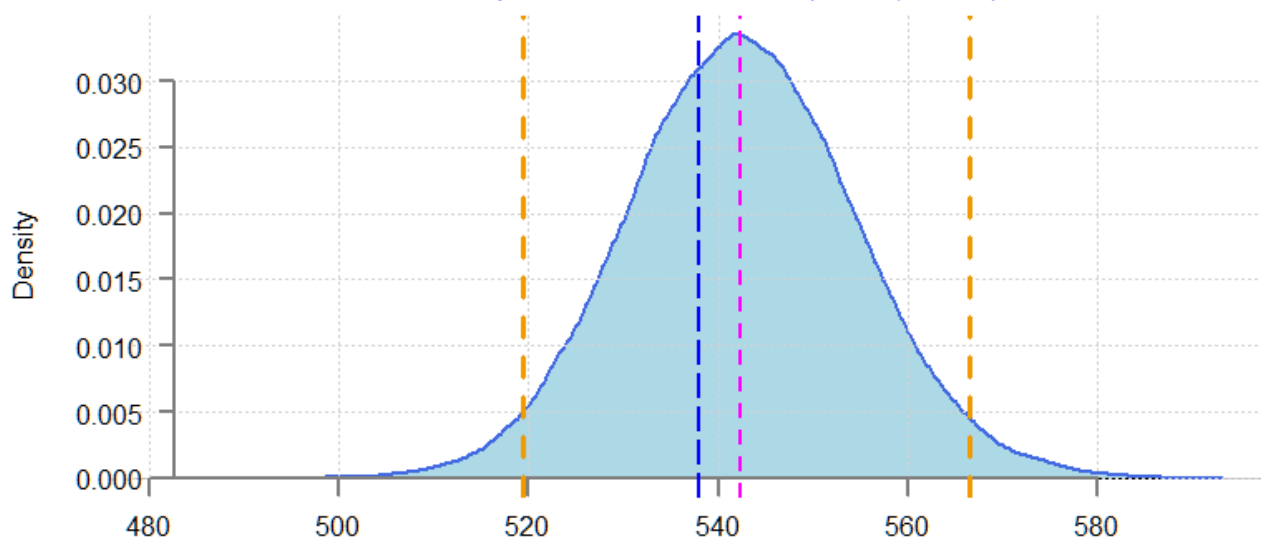
# Hypothesized Population 75th Percentile (Third Quartile) value
qnt0 = 600

# "Metric" is qnt quantile
quantile(rnorm(n,sm,ssd),qnt)

# Get the desired quantile of 50,000 replicates and plot density
m = 50000
sdv3a = replicate(m,quantile(rnorm(n,sm,ssd),qnt))
mysdci(sdv3a,pcol="lightblue",dcol="royalblue",
      main=p("Pop.",100*qnt,"th Percentile Sampling Distr.
            from Normal samples of size ",n,sep=""),
      sub="assuming GMAT data mean and variance")
```

```
abline(v=qnt0,lwd=2,col="red2")      # hypothesized value: qnt0 = 600
abline(v=sqnt,lwd=2,col="blue",lty=5) # GMAT data 75th Empirical Percentile
```

Pop.75th Percentile Sampling Distr. from Normal samples of size 85
 assuming GMAT data mean and variance
 95% Empirical Confidence Interval: (519.33, 566.45)

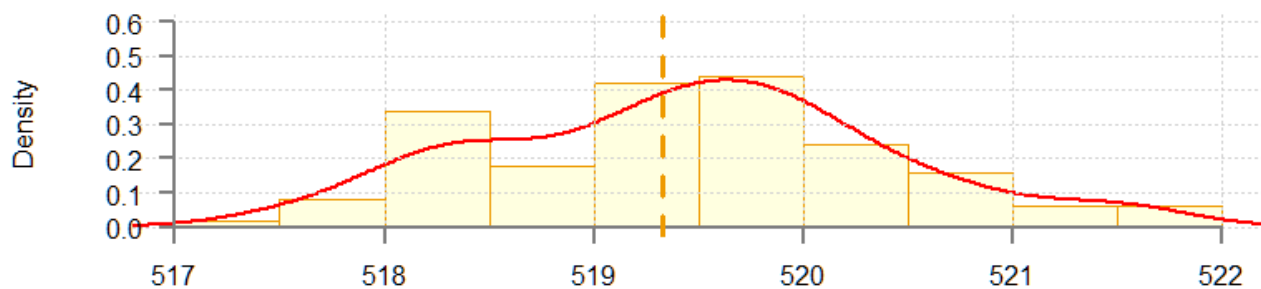


Test conclusion: Since 600 is outside the conf interval we reject the null-hypothesis

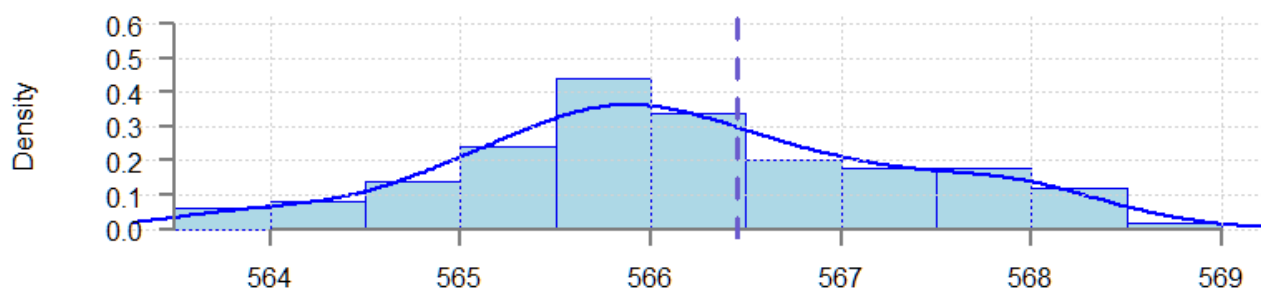
CI variation:

```
m = 1000; M = 100
CIsdv3a = rpl(M,myeCI(rpl(m,quantile(rnorm(n,sm,ssd),qnt))))
myCIsdv(CIsdv3a,myeCI(sdv3a))
myCIsdv(CIsdv3a,myeCI(sdv3a),ylim=c(0,0.6))
```

95% CI's left end-points (519.33 indicated)



95% CI's right end-points (566.45 indicated)



TRY repeating for other percentiles

Second approach: Make NO assumptions about population distribution, instead use only the data points and create simulated samples from these points (re-sampling)

```
# The sampling distr. will be obtained by "resampling": use the built-in function "sample"
x = seq(1,5)

sample(x,5,replace=T)    # we use re-sampling with replacement
sample(x,5,replace=F)

sample(x,12,replace=T)
# sample(x,12,replace=F)
# -----
```

Example 1 via resampling

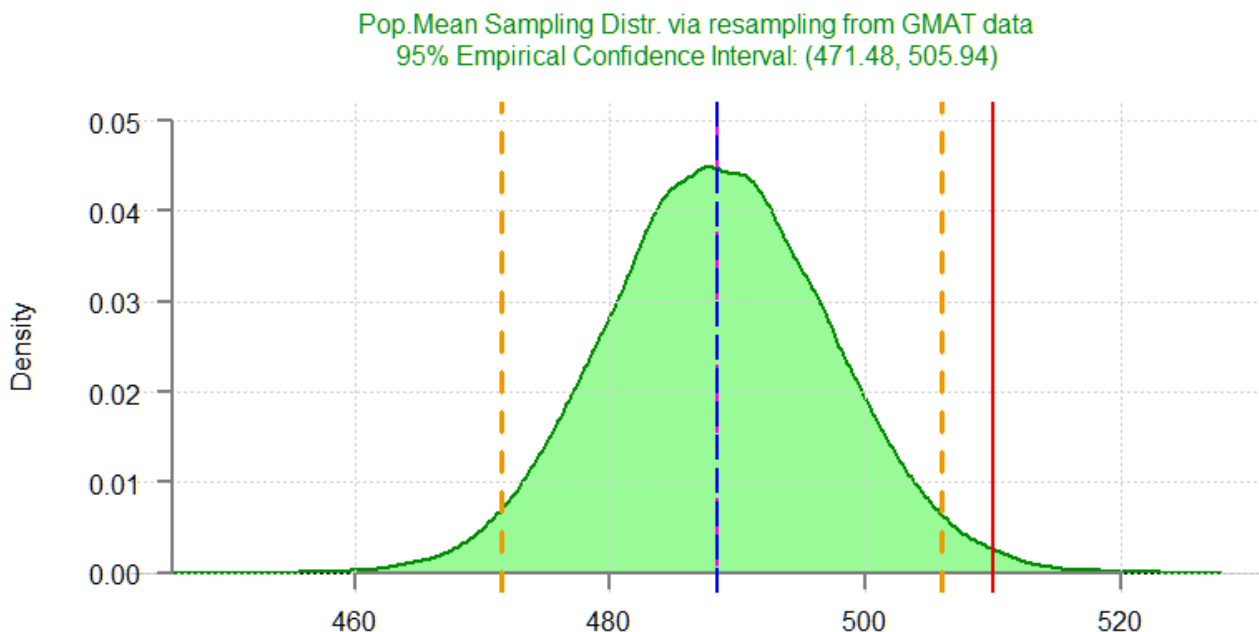
```
# Empirical Confidence Interval for the population mean
# "Null-hypothesis: the population mean of GMAT data is 510"

# Sample size (note: we do not need sm and ssd anymore)
n = length(GMAT); n

# In our case, we re-sample from the GMAT vector
s = sample(GMAT,n,replace=T)

# check the mean of the sample obtained by re-sampling
mean(s)

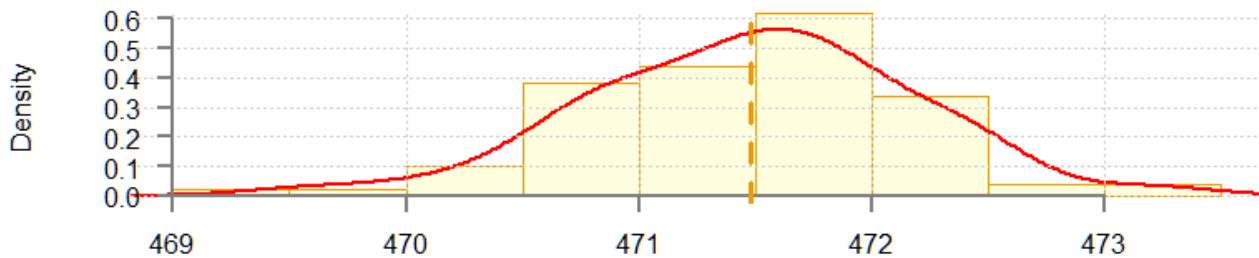
# Replicate it
m = 50000
sdv1b = replicate(m,mean(sample(GMAT,n,T)))
# and plot sampling distribution and CI
mysdci(sdv1b,ylim=c(0,0.05),
        main="Pop.Mean Sampling Distr. via resampling from GMAT data")
abline(v=mu0,lwd=2,col="red2")      # hypothesized value: mu0 = 510
abline(v=sm,lwd=2,col="blue",lty=5) # GMAT data sample mean (average)
```



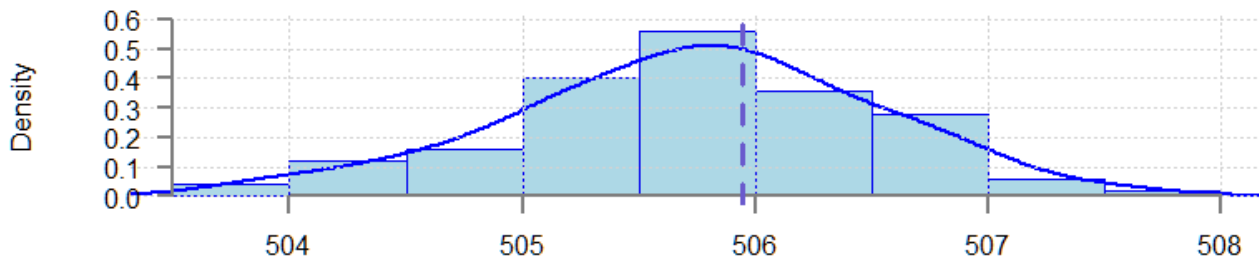
```
# CI variation:
m = 1000; M = 100
CIsdv1b = rpl(M,myeCI(rpl(m,mean(sample(GMAT,n,T)))))
```

```
myCIsdv(CIsdv1b,myeCI(sdv1b))
myCIsdv(CIsdv1b,myeCI(sdv1b),ylim=c(0,0.6))
```

95% CI's left end-points (471.48 indicated)

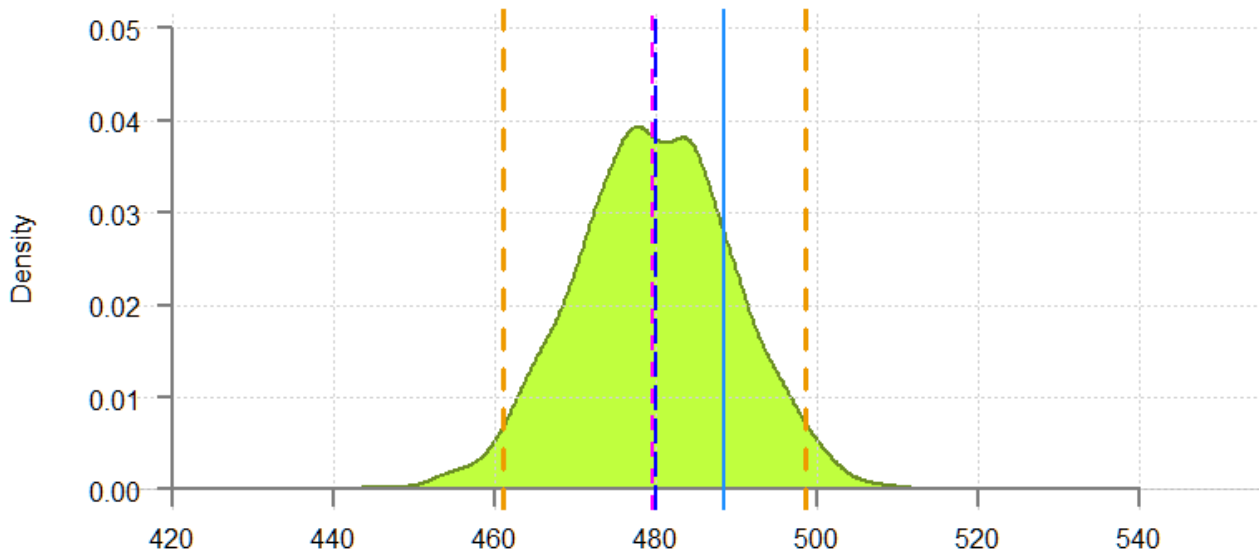


95% CI's right end-points (505.94 indicated)



```
# Try it for another N(sm,ssd^2) vector
w = rnorm(n,sm,ssd)
sdvn = replicate(m,mean(sample(w,n,T)))
mysdci(sdvn,pcol="olivedrab1",dcol="olivedrab",xlim=c(420,550),ylim=c(0,0.05),
        main=p("Pop.Mean Sampling Distr. via resampling from random Normal sample of size",n),
        sub="assuming GMAT data mean and variance")
abline(v=mean(w),lwd=2,col="blue",lty=5)      # w sample mean (average)
abline(v=sm,lwd=2,col="dodgerblue",lty=1)     # GMAT average
```

Pop.Mean Sampling Distr. via resampling from random Normal sample of size 85
assuming GMAT data mean and variance
95% Empirical Confidence Interval: (461.01, 498.55)



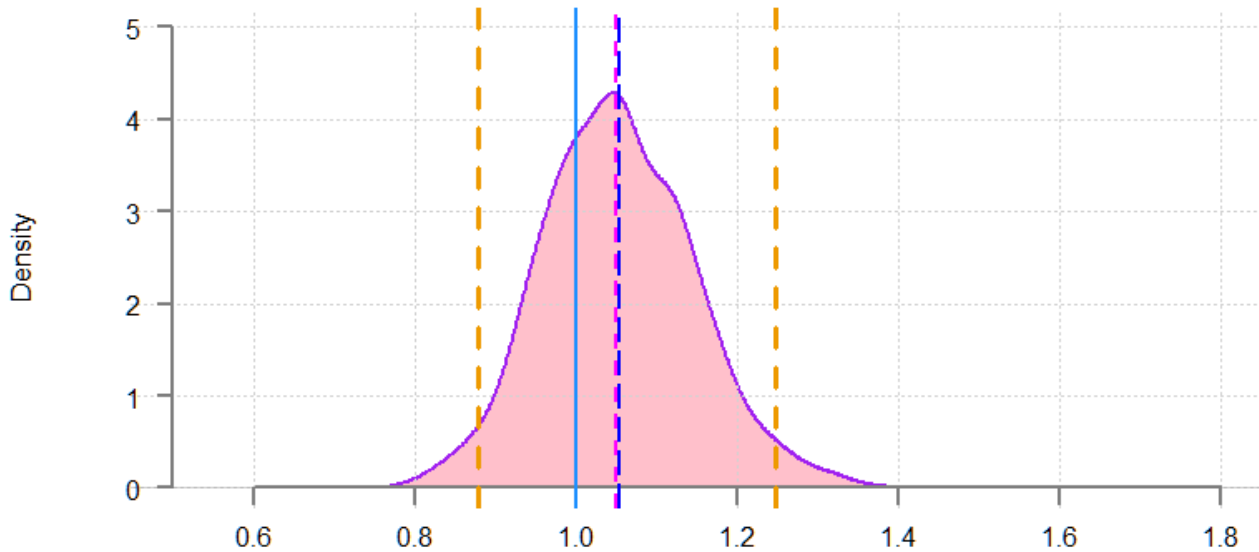
```
# Try it for another population distribution
```

```

w = rexp(n)
sdve = replicate(m,mean(sample(w,n,T)))
mysdci(sdve,pcol="pink",dcol="purple",xlim=c(0.5,1.8),ylim=c(0,5),
        main=p("Pop.Mean Sampling Distr. via resampling from",n,"Exp(1) data pts"))
abline(v=mean(w),lwd=2,col="blue",lty=5)      # w sample mean (average)
abline(v=1,lwd=2,col="dodgerblue",lty=1)      # pop mean of Exp(1)

```

Pop.Mean Sampling Distr. via resampling from 85 Exp(1) data pts
95% Empirical Confidence Interval: (0.88, 1.25)



Example 2 via resampling

```

# Empirical Confidence Interval for the population standard deviation
# "Null-hypothesis: population standard deviation is 78"

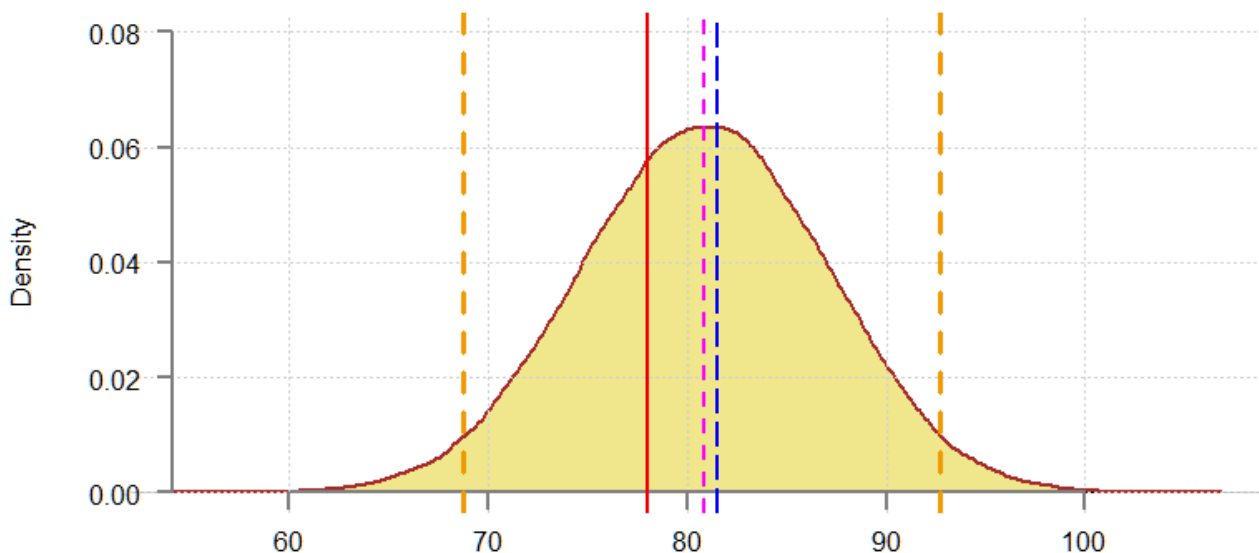
```

```

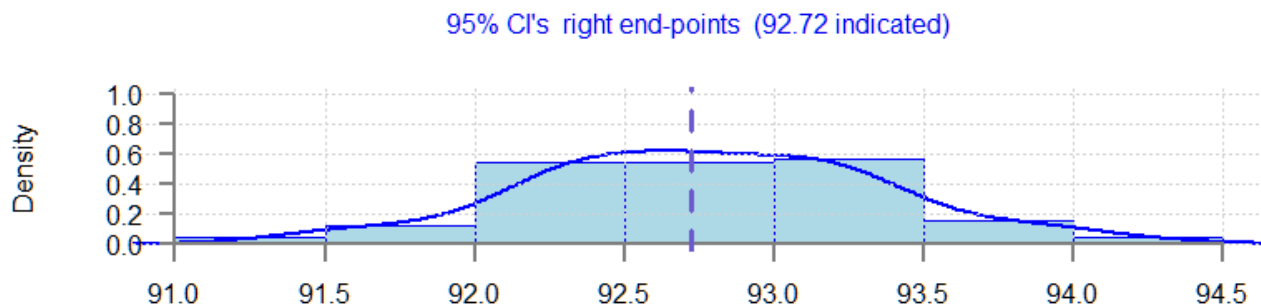
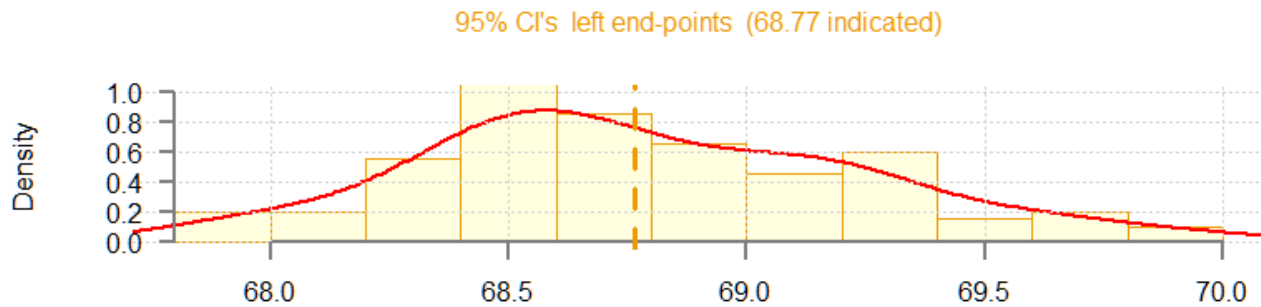
m = 50000
sdv2b = replicate(m,sd(sample(GMAT,n,T)))
mysdci(sdv2b,ylim=c(0,0.08),pcol="khaki",dcol="brown",
        main="Pop.St.Dev. Sampling Distr. via resampling from GMAT data")
abline(v=sd0,lwd=2,col="red2")      # hypothesized value: sd0 = 78
abline(v=ssd,lwd=2,col="blue",lty=5) # GMAT data Sample Standard Deviation

```

Pop.St.Dev. Sampling Distr. via resampling from GMAT data
95% Empirical Confidence Interval: (68.77, 92.72)



```
# CI variation:
m = 1000; M = 100
CIsdv2b = repl(M,myeCI(repl(m,sd(sample(GMAT,n,T))))))
myCIsdv(CIsdv2b,myeCI(sdv2b))
myCIsdv(CIsdv2b,myeCI(sdv2b),ylim=c(0,1))
```

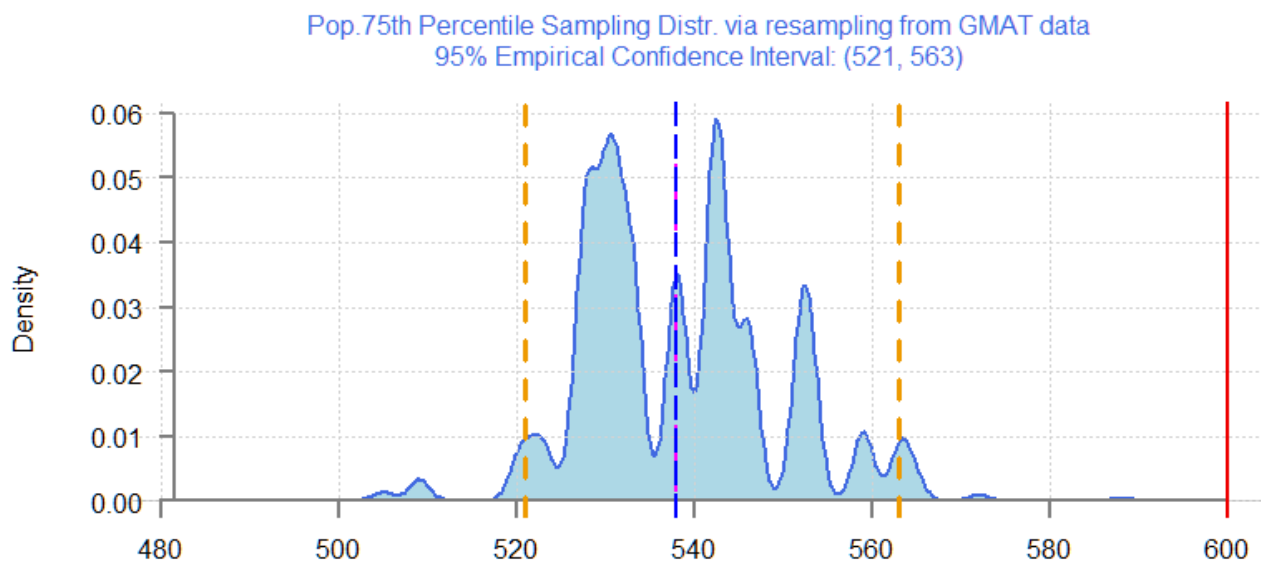


Example 3 via resampling

```
# Empirical Confidence Interval for the population 75th percentile
# "Null-hypothesis: 75th percentile of the population is 600"

qnt = 0.75          # qnt = 0.5
m = 50000

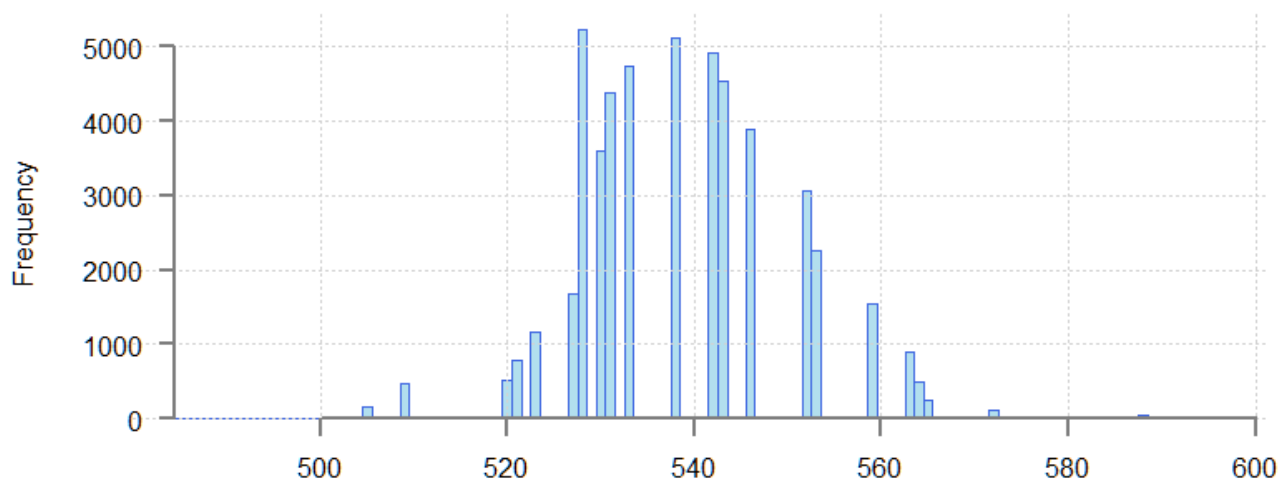
sdv3b = replicate(m,quantile(sample(GMAT,n,T),qnt))
# auxiliary string for the super-long chart title
main = p("Pop.",100*qnt,
         "th Percentile Sampling Distr. via resampling from GMAT data",sep="")
mysdci(sdv3b,pcol="lightblue",dcol="royalblue",main=main)
abline(v=qnt0,lwd=2,col="red2")      # hypothesized value: qnt0 = 600
abline(v=sqnt,lwd=2,col="blue",lty=5) # GMAT data 75th Empirical Percentile
```



```
# Why such bumpiness?
summary(sdv3b)
# Key: How many distinct values does vector 'sampldist' have?
length(sdv3b); length(unique(sdv3b))

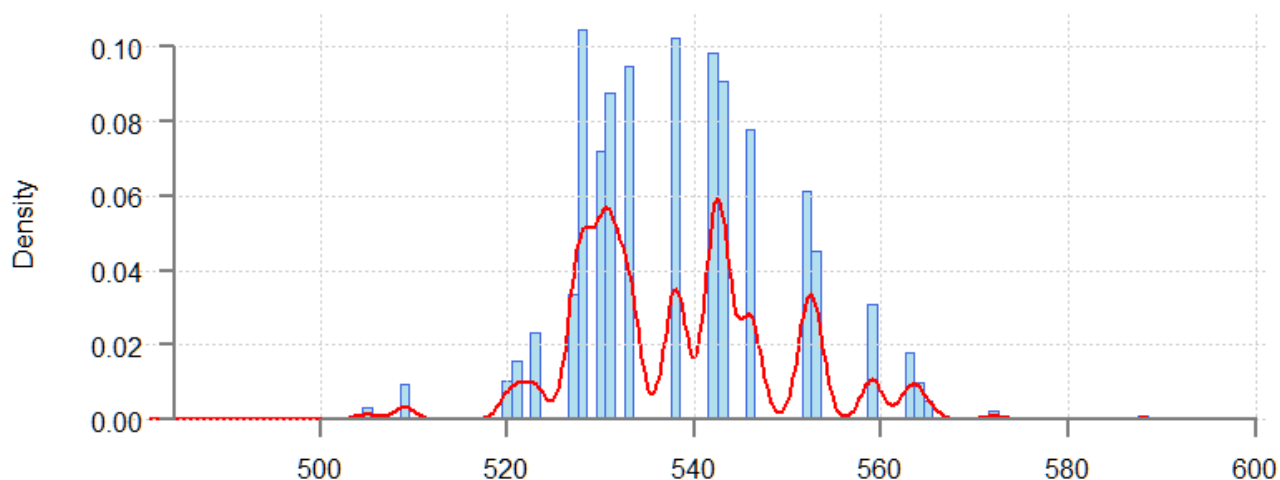
# Empirical density does a poor job in representing sparse data; try histogram
myhist(sdv3b,freq=T,main=main,col="lightblue2",border="royalblue")
myhist(sdv3b,freq=T,breaks=seq(min(sdv3b)-0.5,max(sdv3b)+0.5),
      main=main,col="lightblue2",border="royalblue")
```

Pop.75th Percentile Sampling Distr. via resampling from GMAT data: 50,000 data points



```
myhd(sdv3b,breaks=seq(min(sdv3b)-0.5,max(sdv3b)+0.5),cntFlg=F,main=main,
     col="lightblue2",border="royalblue")
```

Pop.75th Percentile Sampling Distr. via resampling from GMAT data

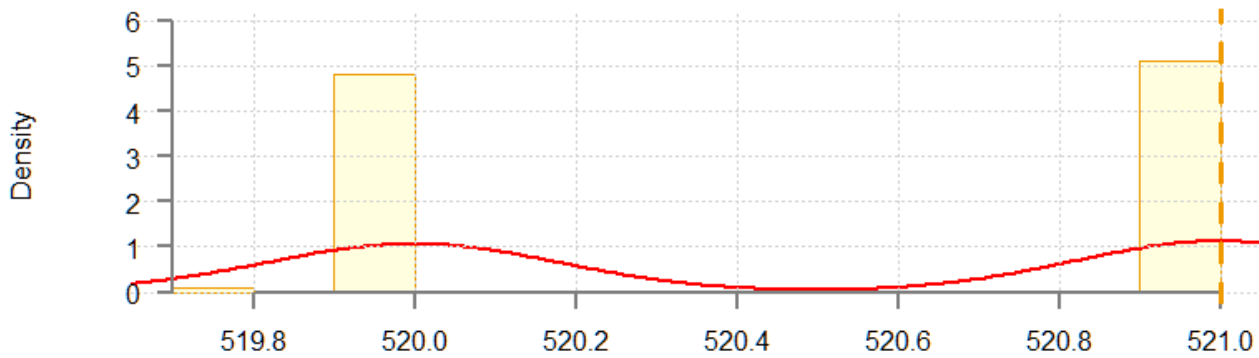


```
myhd2(sdv3b,cntFlg=F,main=main,col="lightblue2",border="royalblue")
```

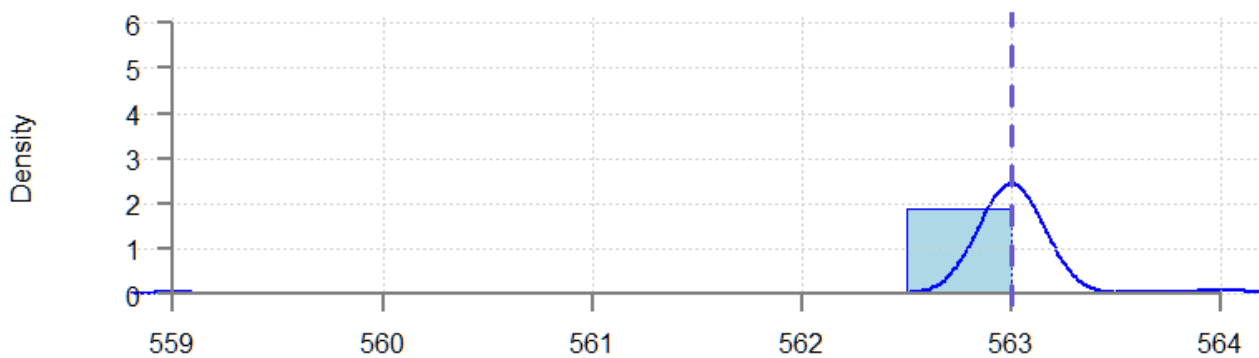
sampling distribution for percentiles is very bumpy; What about CIs?

```
# CI variation:
m = 1000; M = 100
CIsv3b = rpl(M,myeCI(rpl(m,quantile(sample(GMAT,n,T),qnt))))
myCIsdv(CIsv3b,myeCI(sdv3b))
myCIsdv(CIsv3b,myeCI(sdv3b),ylim=c(0,6))
```

95% CI's left end-points (521 indicated)



95% CI's right end-points (563 indicated)



Unusual distribution? Check the actual frequencies:

`table(CIsdv3b[1,])` # Frequency table of left end-points

519.725	520	520.975	521
1	48	4	47

`table(CIsdv3b[2,])` # Frequency table of right end-points

559	563	563.025	564
1	95	1	3

Recall: sampling distribution for percentiles is very bumpy: this will be the case for all percentiles: try median (uncomment "qnt = 0.5" line)

Try it for a Normal $N(0,1)$ sample of size 85 = `length(GMAT)`

`n = length(GMAT)`

`n = 5000`

`w = rnorm(n)`

`qnt = 0.75`

`m = 50000`

`sdvnr = replicate(m, quantile(sample(w,n,T),qnt))`

`summary(sdvnr)`

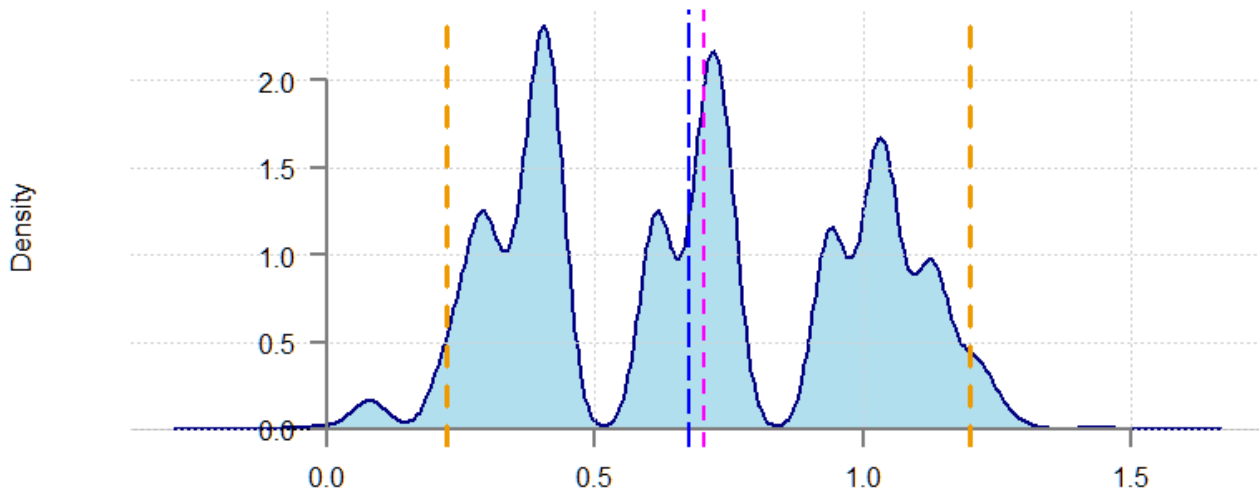
`length(sdvnr); length(unique(sdvnr))`

`main = p("Pop.", 100*qnt,`
`"th Percentile Sampling Distr. via resampling from $N(0,1)$ sample of size ", n, sep="")`

`mysdci(sdvnr, pcol="lightblue2", dcol="navy", main=main)`

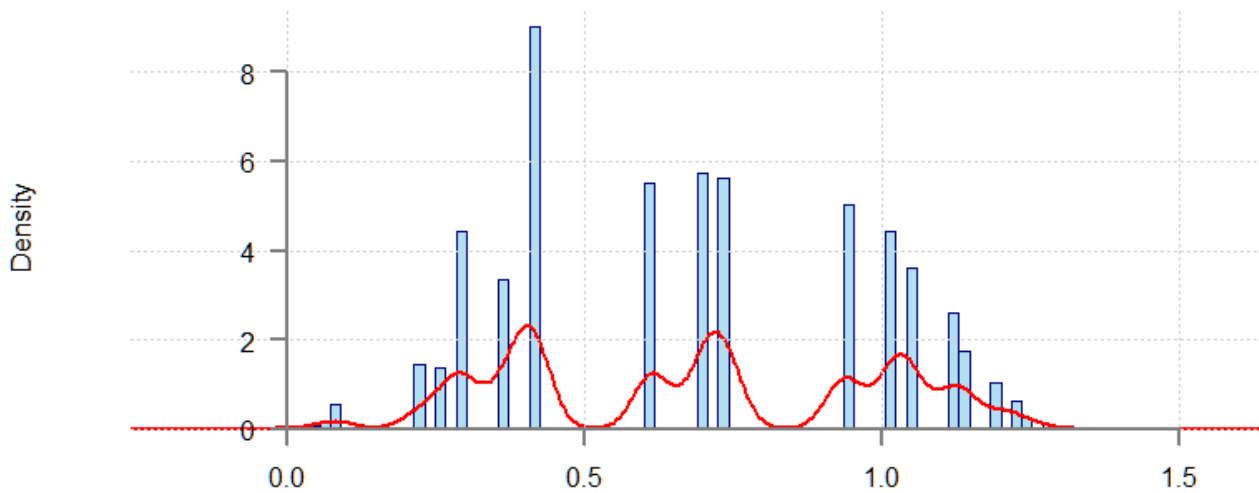
`abline(v=qnorm(qnt), lwd=2, col="blue", lty=5)` # Pop. 75th Percentile of $N(0,1)$

Pop.75th Percentile Sampling Distr. via resampling from $N(0,1)$ sample of size 85
95% Empirical Confidence Interval: (0.22, 1.2)



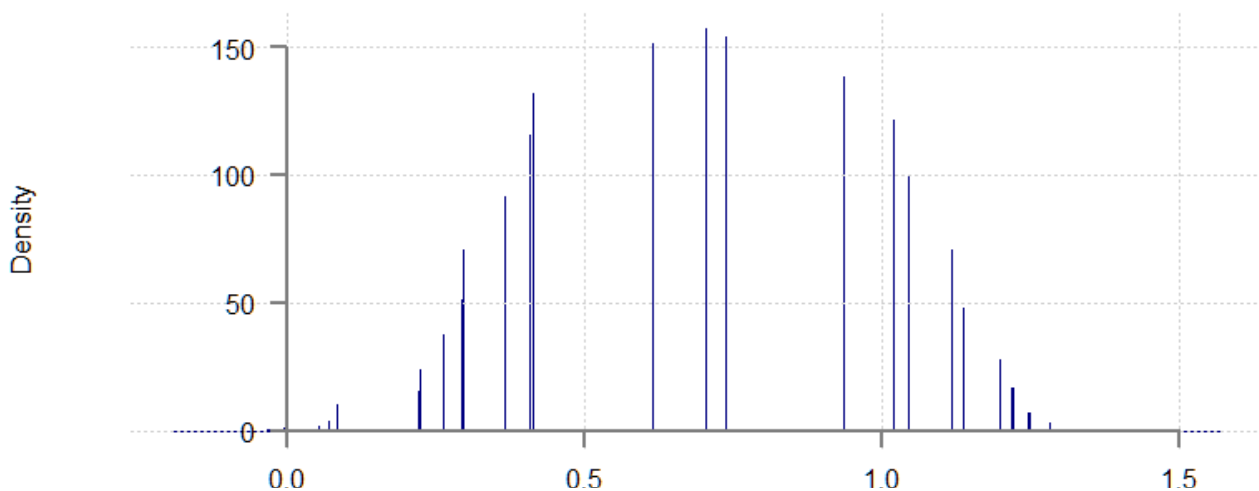
```
myhd(sdvnr,breaks=seq(min(sdvnr),max(sdvnr),l=101),cntFlg=F,
     main=main,col="lightblue2",border="navy")
```

Pop.75th Percentile Sampling Distr. via resampling from $N(0,1)$ sample of size 85



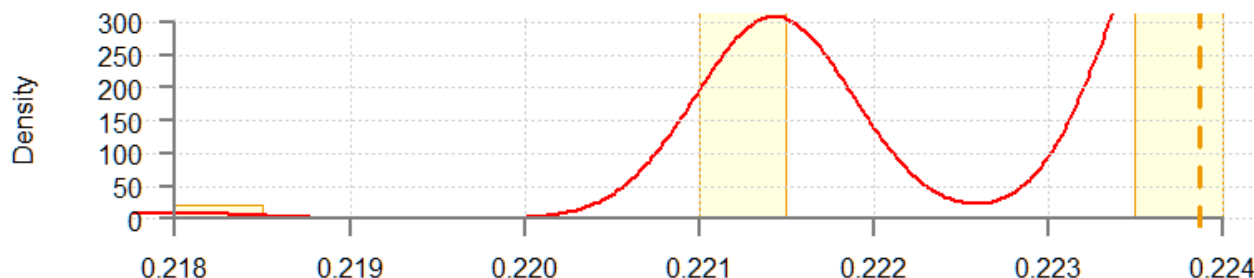
```
myhd2(sdvnr,cntFlg=F,main=main,col="lightblue2",border="navy")
```

Pop.75th Percentile Sampling Distr. via resampling from $N(0,1)$ sample of size 85
(discrete data)

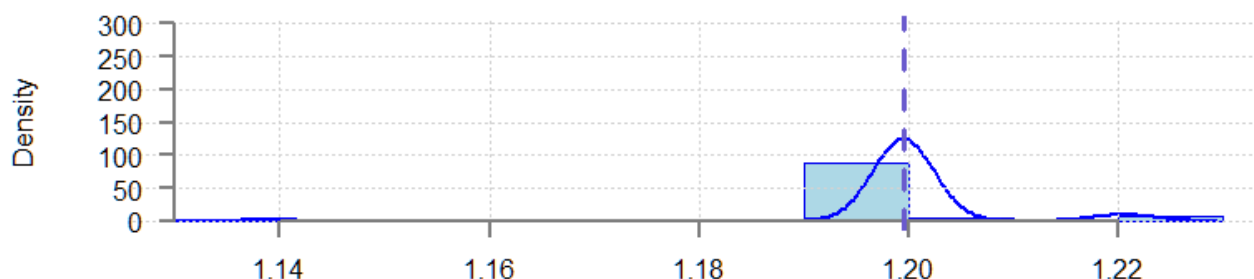


```
# CI variation:
m = 1000; M = 100
CIsdvnr = rpl(M,myeCI(rpl(m,quantile(sample(w,n,T),qnt))))
myCIsdv(CIsdvnr,myeCI(sdvnr))
```

95% CI's left end-points (0.22 indicated)



95% CI's right end-points (1.2 indicated)



```
table(CIsdvnr[1,]) # Frequency table of left end-points
```

```
0.218016745973082 0.221430441571721 0.223802533914368 0.223863356794949
1 35 6 58
```

```
table(CIsdvnr[2,]) # Frequency table of right end-points
```

```
1.1388181781256 1.19948646428119 1.2000114595014 1.2204862730898
1 88 4 7
```

```
# Same problems for larger samples (uncomment "n = 5000" line)
```

```
# Conclusion: Bumpier sampling distribution plots, yet confidence intervals are 'stable'
```

```
# -----
```

Example 4:

```
# Can we check if the GPA and GMAT scores from the admission data set are independent?
#
```

```
# The question really is: are the underlying random variables, GPA and GMAT scores of
# the student population, independent?
```

```
#
```

```
# Independence of random variables cannot be checked by calculation.
```

```
# Furthermore, we only have a sample of pairs (GPA, GMAT) of size n = 85.
```

```
#
```

```
# But we can check un-correlated-ness!
```

```
# Correlation: "a quantity measuring the extent of interdependence"
```

```
#
```

```
# Independent random variables are uncorrelated (correlation = 0).
```

```
# The converse is not true: There exist dependent random variables whose correlation is zero.
```

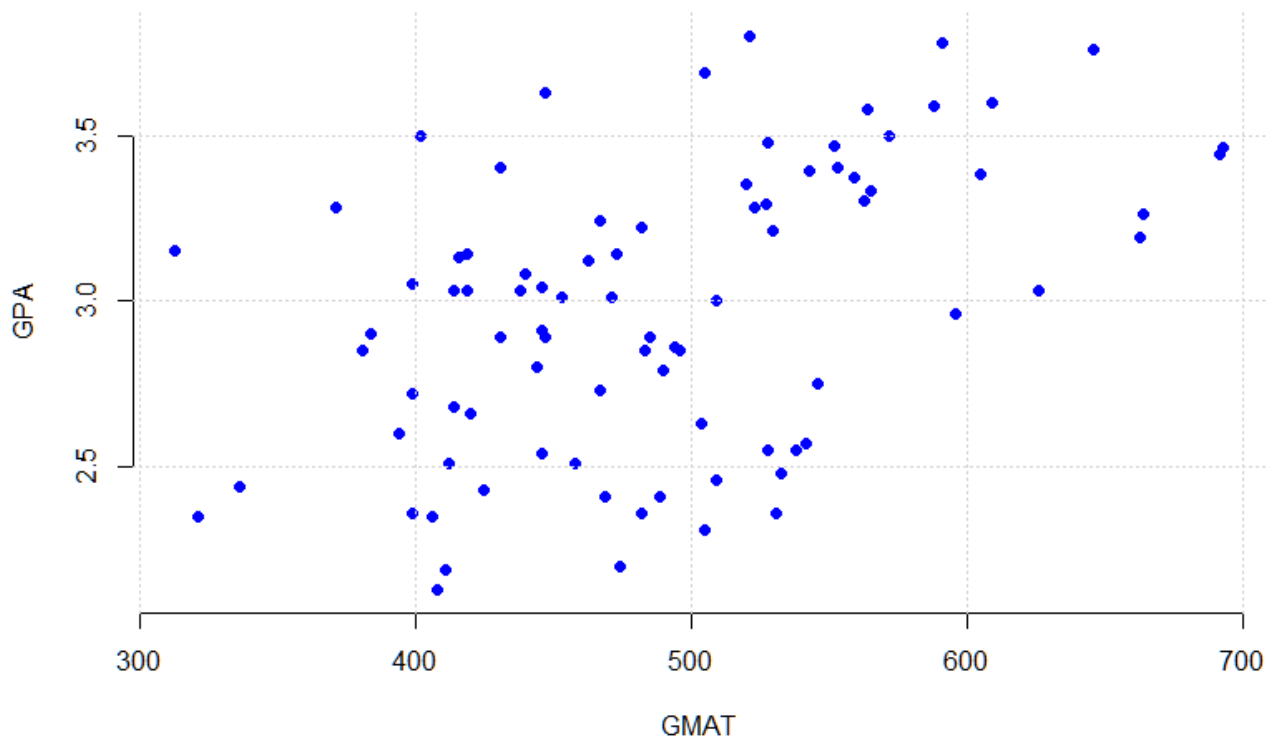
```
# We test the following:
```

```
# Null-hypothesis: GPA and GMAT scores are uncorrelated!

# If the null-hypothesis is accepted, i.e., the GPA and GMAT scores are uncorrelated,
# they may or may not be independent.
# However, if the null hypothesis is rejected in favor of the alternative:
# GPA and GMAT scores are not uncorrelated, hence not independent!
```

```
GPA = adm$GPA
```

```
# Scatter plot the GPA scores against the GMAT scores
plot(GMAT,GPA,pch=19,col="blue",frame.plot=F); grid()
```



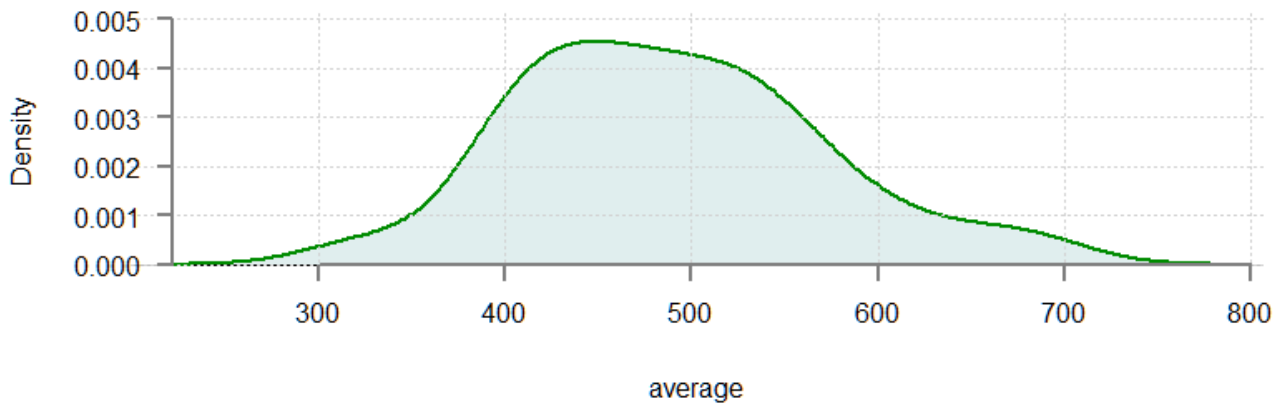
```
# What is the sample correlation?
Scor = cor(GPA,GMAT); Scor
```

```
[1] 0.4606332
```

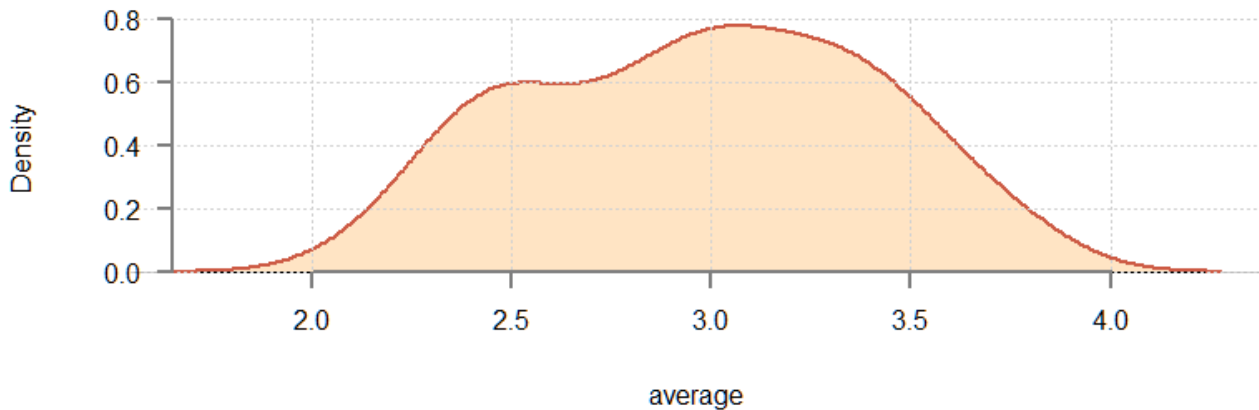
```
sm1 = mean(GPA); ssd1 = sd(GPA)
sm2 = mean(GMAT); ssd2 = sd(GMAT)
```

```
# the empirical densities
par(mfrow=c(2,1)) # set 2x1 plots
myed(GMAT,ylim=c(0,0.005),pcol="azure2",dcol="green4",
     xlab="average",main="GMAT data density")
myed(GPA,pcol="bisque",dcol="coral3",
     xlab="average",main="GPA data density")
par(mfrow=c(1,1)) # set 1x1 plot
```

GMAT data density: 85 data points



GMAT data density: 85 data points



First approach: Assume normality of GPA and GMAT data

```
# Find the confidence interval for the correlation of two 'independent' vectors
# from normal distributions N(sm1,ssd1) and N(sm2,ssd2), respectively, where
# sm1, ssd1, sm2, and ssd2 are sample variance and sample standard deviation
# of GPA and GMAT vectors from the admission data set.

# Question: if GPA and GMAT vectors were independent (more precisely, uncorrelated),
#           how likely it is that their correlation is approx. 0.46?
# We can answer this if we know the distribution of the 'sample' correlation of
# two independent normal random vectors of same length as GPA and GMAT vectors

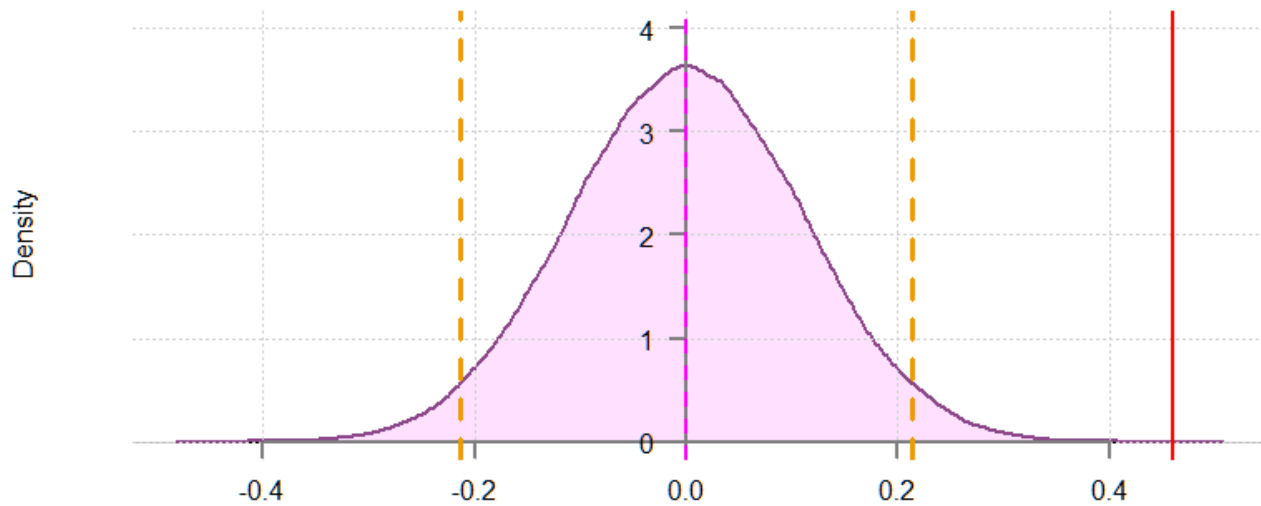
n = length(GPA)
cor(rnorm(n,sm1,ssd1),rnorm(n,sm2,ssd2))

# The sampling distribution is created by replicating this 100 thousand times
m = 100000
sdv4a = replicate(m,cor(rnorm(n,sm1,ssd1),rnorm(n,sm2,ssd2)))

mysdci(sdv4a,ylim=c(0,4),pcol="thistle1",dcol="orchid4",
       main=paste("Pop.Correlation Sampling Distr. from 2 indep Normal samples of size",n),
       sub="assuming GPA / GMAT data mean and variance")

# Where does the correlation of our GPA and GMAT vectors (0.46) fall in?
abline(v=Scor,lwd=2,col="red")
```

Pop. Correlation Sampling Distr. from 2 indep Normal samples of size 85
 assuming GPA / GMAT data mean and variance
 95% Empirical Confidence Interval: (-0.21, 0.21)

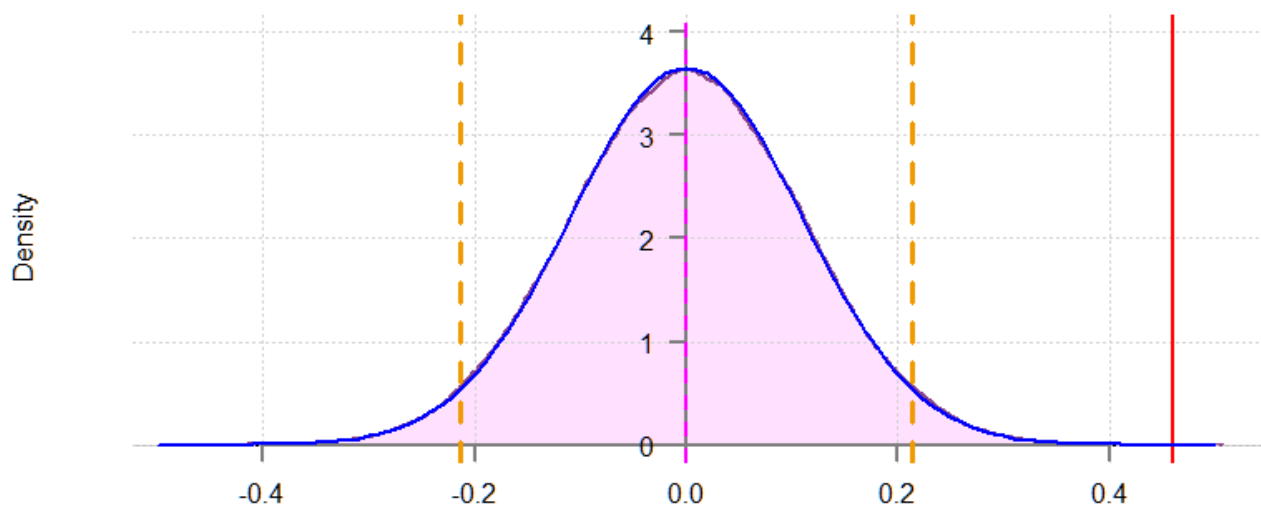


Extremely far in the tail. Check the maximum of the sampling distribution vector
`max(sdv4a)`

Do you accept or reject H_0 : GPA and GMAT scores are independent?

Related question: how close is the sampling distr. to the normal distribution:
`x = seq(-0.5, 0.5, 0.01)`
`lines(x, dnorm(x, 0, sd(sdv4a)), lwd=2, col="blue")`

Pop. Correlation Sampling Distr. from 2 indep Normal samples of size 85
 assuming GPA / GMAT data mean and variance
 95% Empirical Confidence Interval: (-0.21, 0.21)



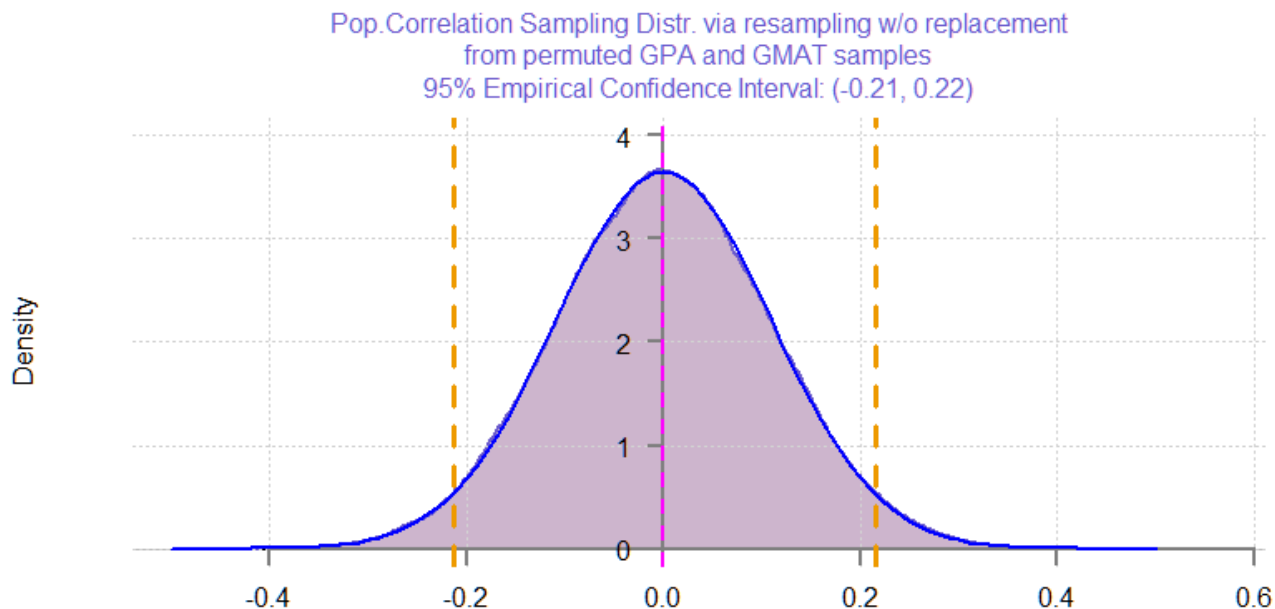
Second approach: Resampling

Recall: resampling (a full length of vector) without replacement is same as permuting it
`v = seq(6)`
`sample(v, 6, replace=FALSE)`

Resampling without replacement (permutations of data vectors)
`m = 100000`
`sdv4b = rpl(m, cor(sample(GPA, n, F), sample(GMAT, n, F)))`

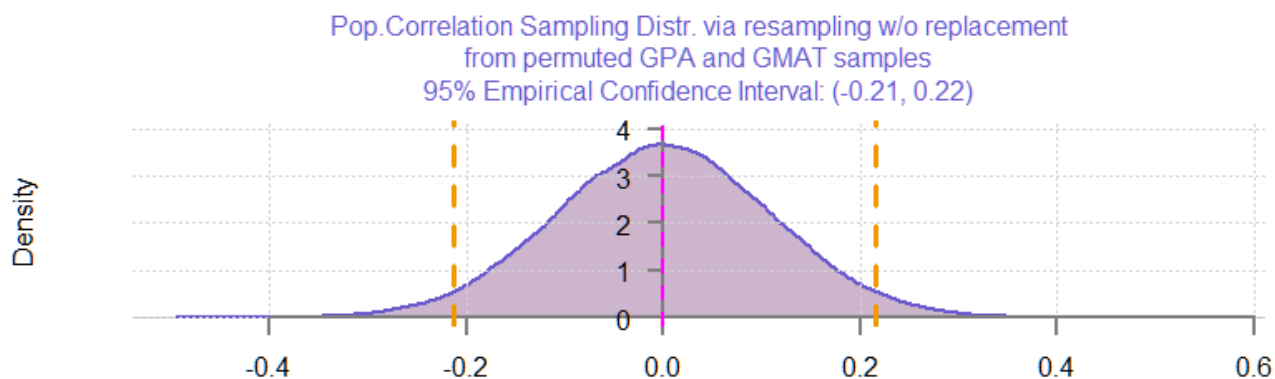
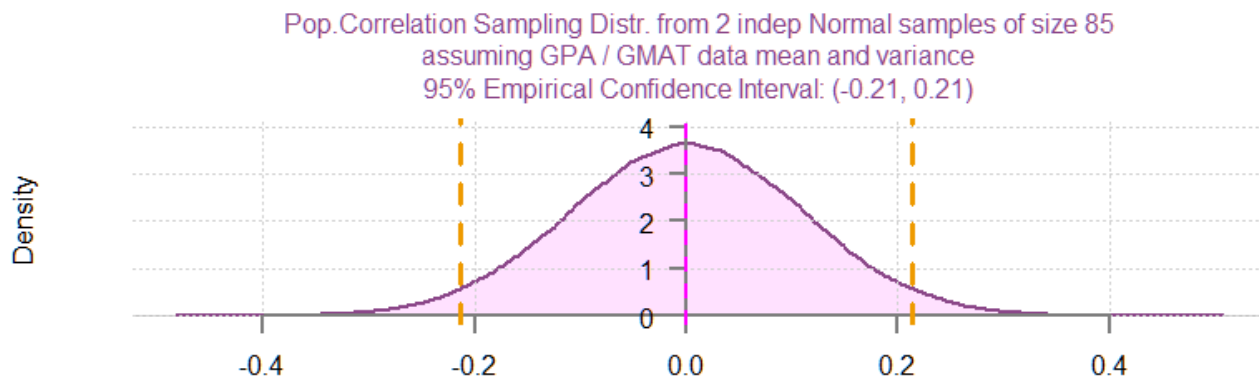
It is not hard to see that it is enough to permute only one of the vectors

```
sdv4b = rpl(m,cor(GPA,sample(GMAT,n,F)))
mysdci(sdv4b,ylim=c(0,4),pcol="thistle3",dcol="slateblue",
      main=paste("Pop.Correlation Sampling Distr. via resampling w/o replacement"),
      sub="from permuted GPA and GMAT samples")
x = seq(-0.5,0.5,0.01)
lines(x,dnorm(x,0,sd(sdv4b)),lwd=2,col="blue")
```



Do you accept or reject H_0 : GPA and GMAT scores are independent?

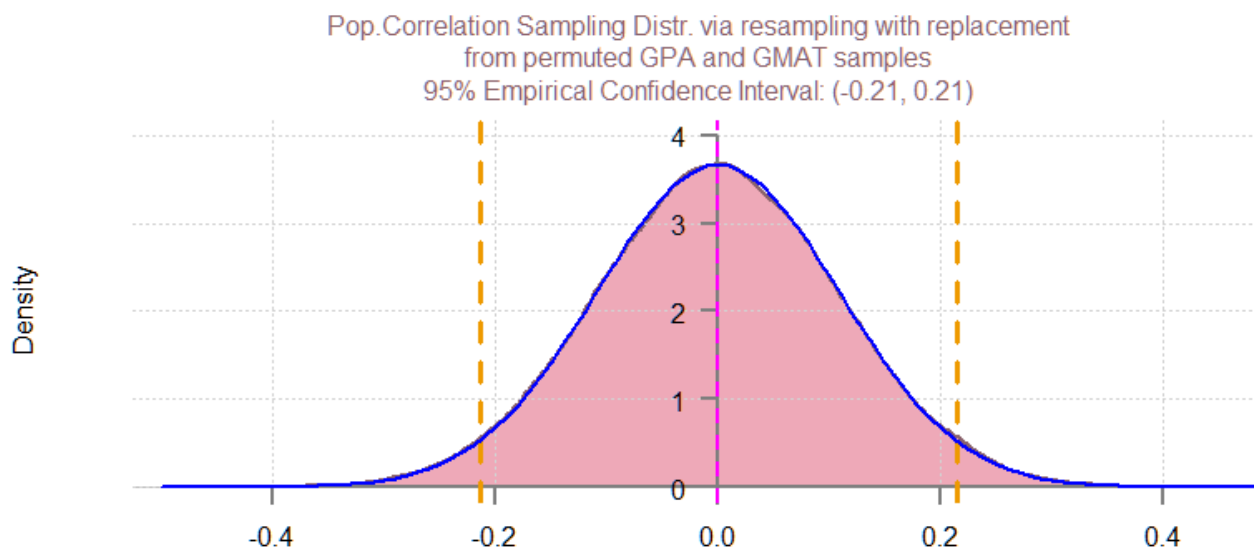
```
# Plot both sample distributions from 4a and 4b
par(mfrow=c(2,1)) # set 2x1 plots
mysdci(sdv4a,ylim=c(0,4),pcol="thistle1",dcol="orchid4",
      main=paste("Pop.Correlation Sampling Distr. from 2 indep Normal samples of size",n),
      sub="assuming GPA / GMAT data mean and variance")
mysdci(sdv4b,ylim=c(0,4),pcol="thistle3",dcol="slateblue",
      main=paste("Pop.Correlation Sampling Distr. via resampling w/o replacement"),
      sub="from permuted GPA and GMAT samples")
par(mfrow=c(1,1)) # set 1x1 plot
```



```
# Interesting observations:
```

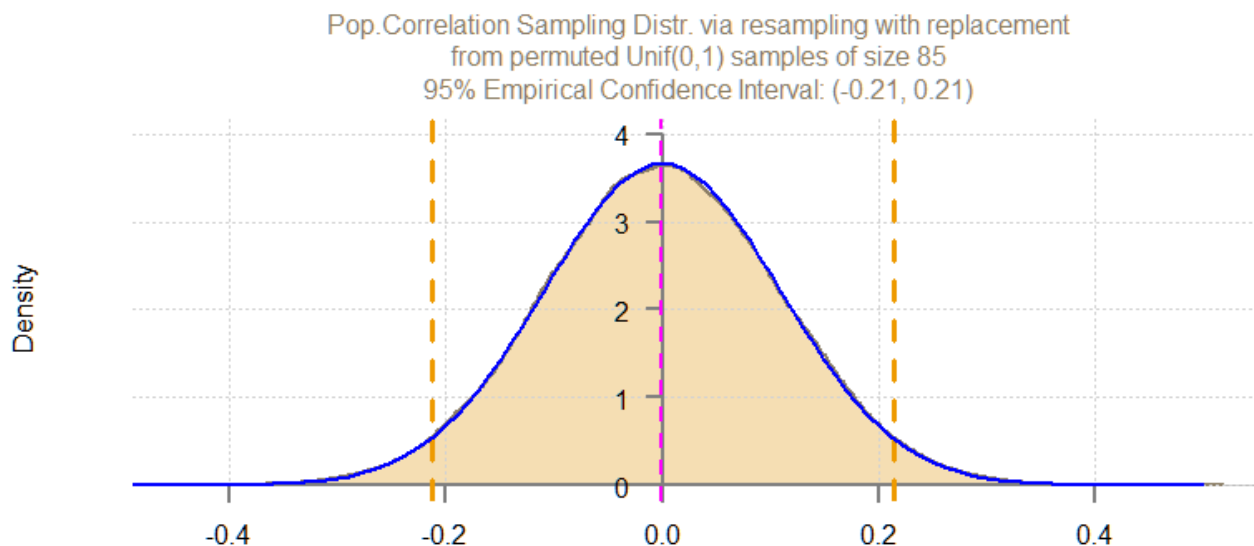
```
# Try resampling WITH replacement - the sampling distr. look pretty much the same
```

```
sdv4c = rpl(m,cor(sample(GPA,n,T),sample(GMAT,n,T)))
mysdci(sdv4c,ylim=c(0,4),pcol="pink2",dcol="pink4",
      main=paste("Pop.Correlation Sampling Distr. via resampling with replacement"),
      sub="from permuted GPA and GMAT samples")
x = seq(-0.5,0.5,0.01)
lines(x,dnorm(x,0,sd(sdv4c)),lwd=2,col="blue")
```

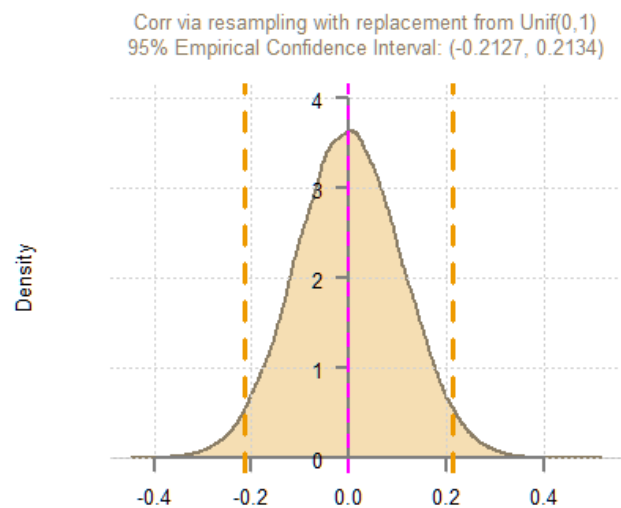
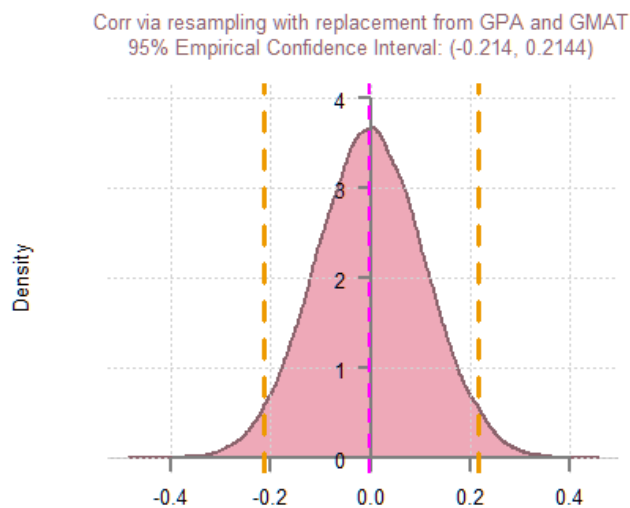
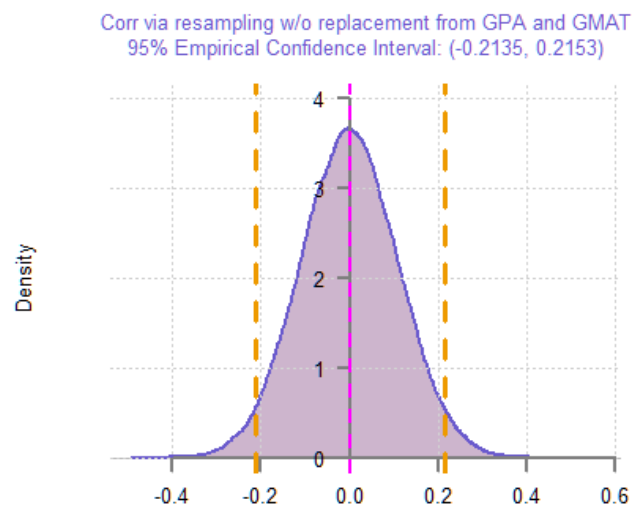
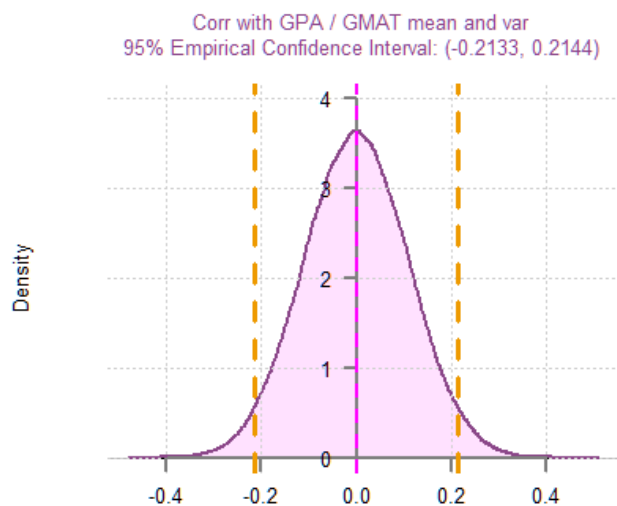


```
# In fact, replace the GPA and GMAT vectors by any two 'independent' generated vectors  
# of length 85. The sampling distribution will look identical
```

```
sdv4d = rpl(m,cor(runif(n),runif(n)))
mysdci(sdv4d,ylim=c(0,4),pcol="wheat",dcol="wheat4",
      main=paste("Pop.Correlation Sampling Distr. via resampling with replacement"),
      sub=paste("from permuted Unif(0,1) samples of size",n))
x = seq(-0.5,0.5,0.01)
lines(x,dnorm(x,0,sd(sdv4d)),lwd=2,col="blue")
```



```
# All-star lineup
par(mfrow=c(2,2))                                # set 2x2 plots
mysdci(sdv4a,ylim=c(0,4),pcol="thistle1",dcol="orchid4",rnd=4,
      main="Corr with GPA / GMAT mean and var")
mysdci(sdv4b,ylim=c(0,4),pcol="thistle3",dcol="slateblue",rnd=4,
      main="Corr via resampling w/o replacement from GPA and GMAT")
mysdci(sdv4c,ylim=c(0,4),pcol="pink2",dcol="pink4",rnd=4,
      main="Corr via resampling with replacement from GPA and GMAT")
mysdci(sdv4d,ylim=c(0,4),pcol="wheat",dcol="wheat4",rnd=4,
      main="Corr via resampling with replacement from Unif(0,1)")
par(mfrow=c(1,1))                                # set 1x1 plot
```



```
# Hence the distributions of the vectors play no role here - only the fact that
# the vectors were independently generated (via R random number generator: the
# vectors are just uncorrelated) matters in shaping the sampling distribution.
```